

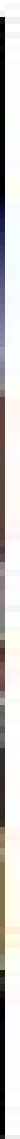
Seismic Imaging on NVIDIA GPUs

**Scott Morton
Hess Corporation**

- Seismic data & imaging
- NVIDIA GPUs + CUDA
 - Why?
 - How?
- Three imaging methods
 - Algorithm
 - Challenges
 - Performance

Seismic Imaging

Data

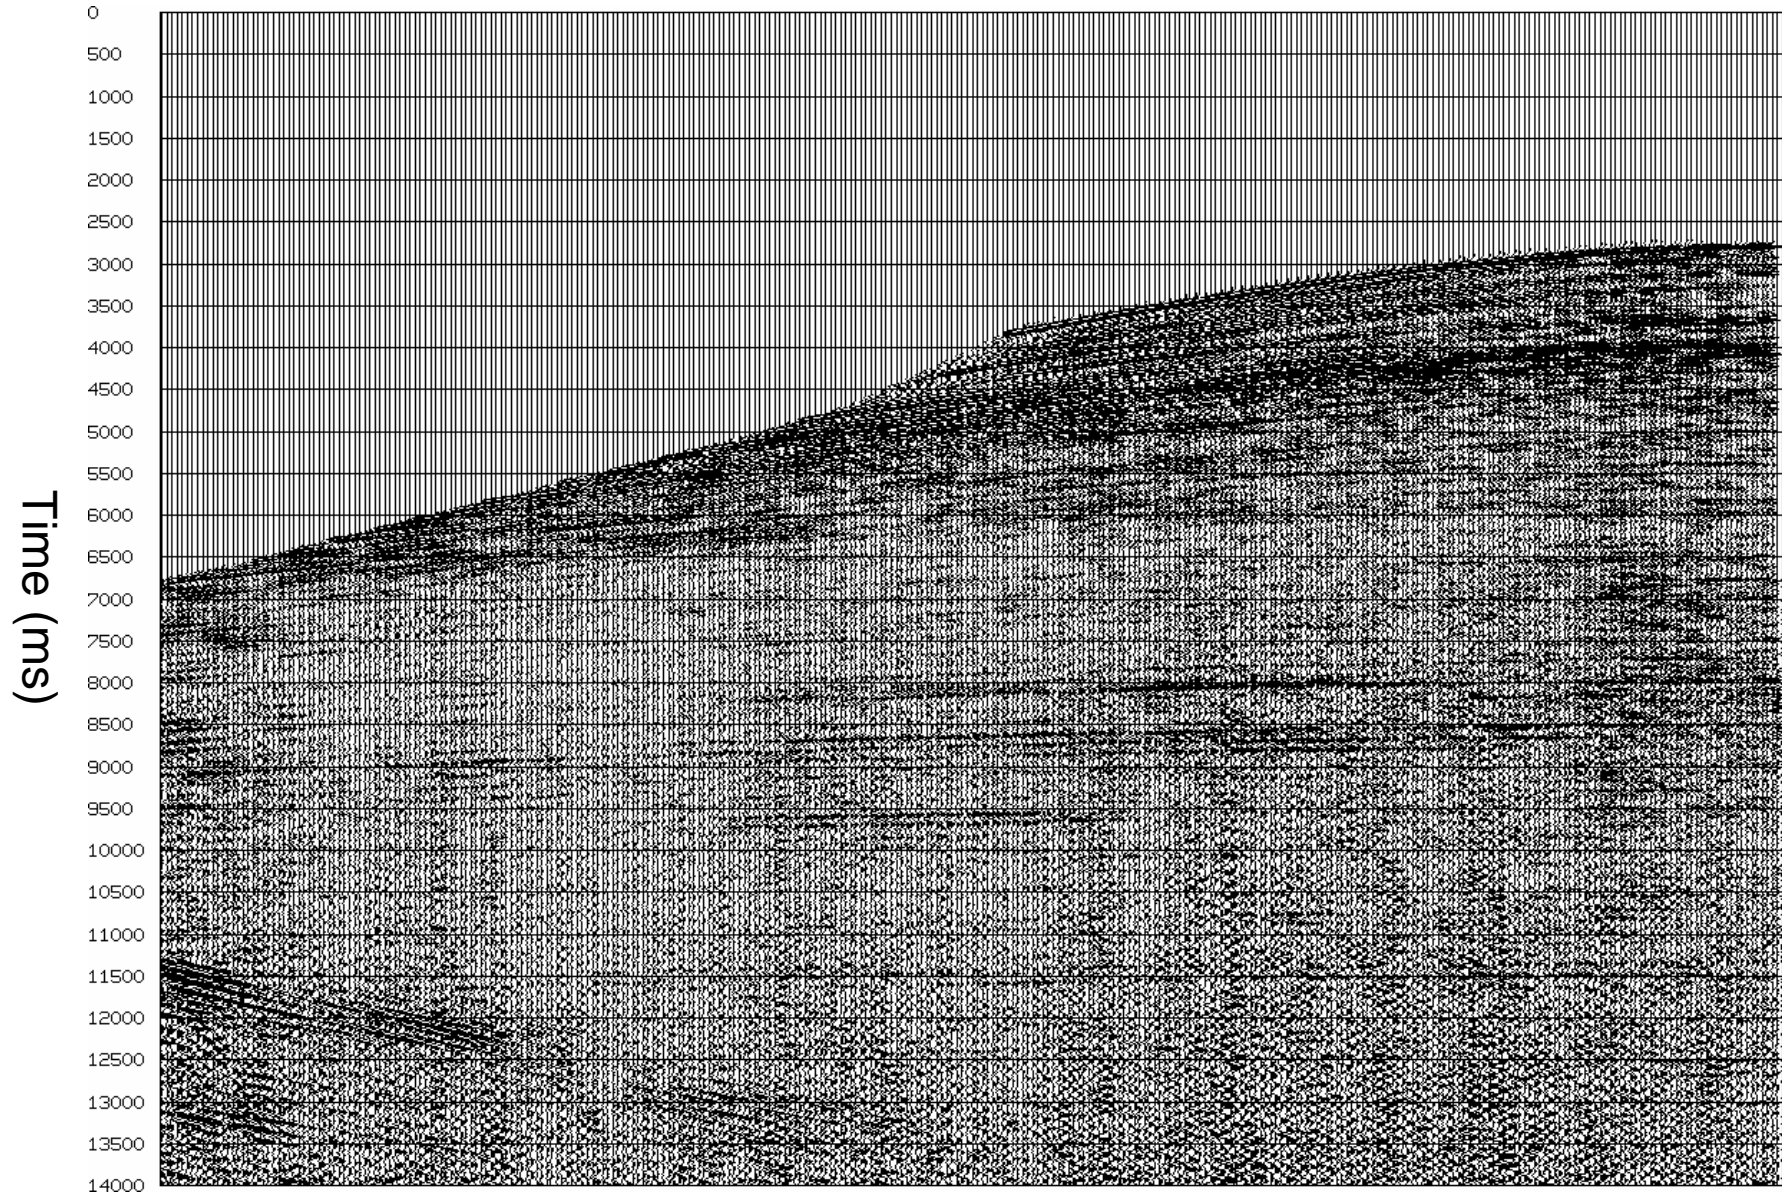


Seismic Imaging

Data



Receiver



Seismic Imaging

Data

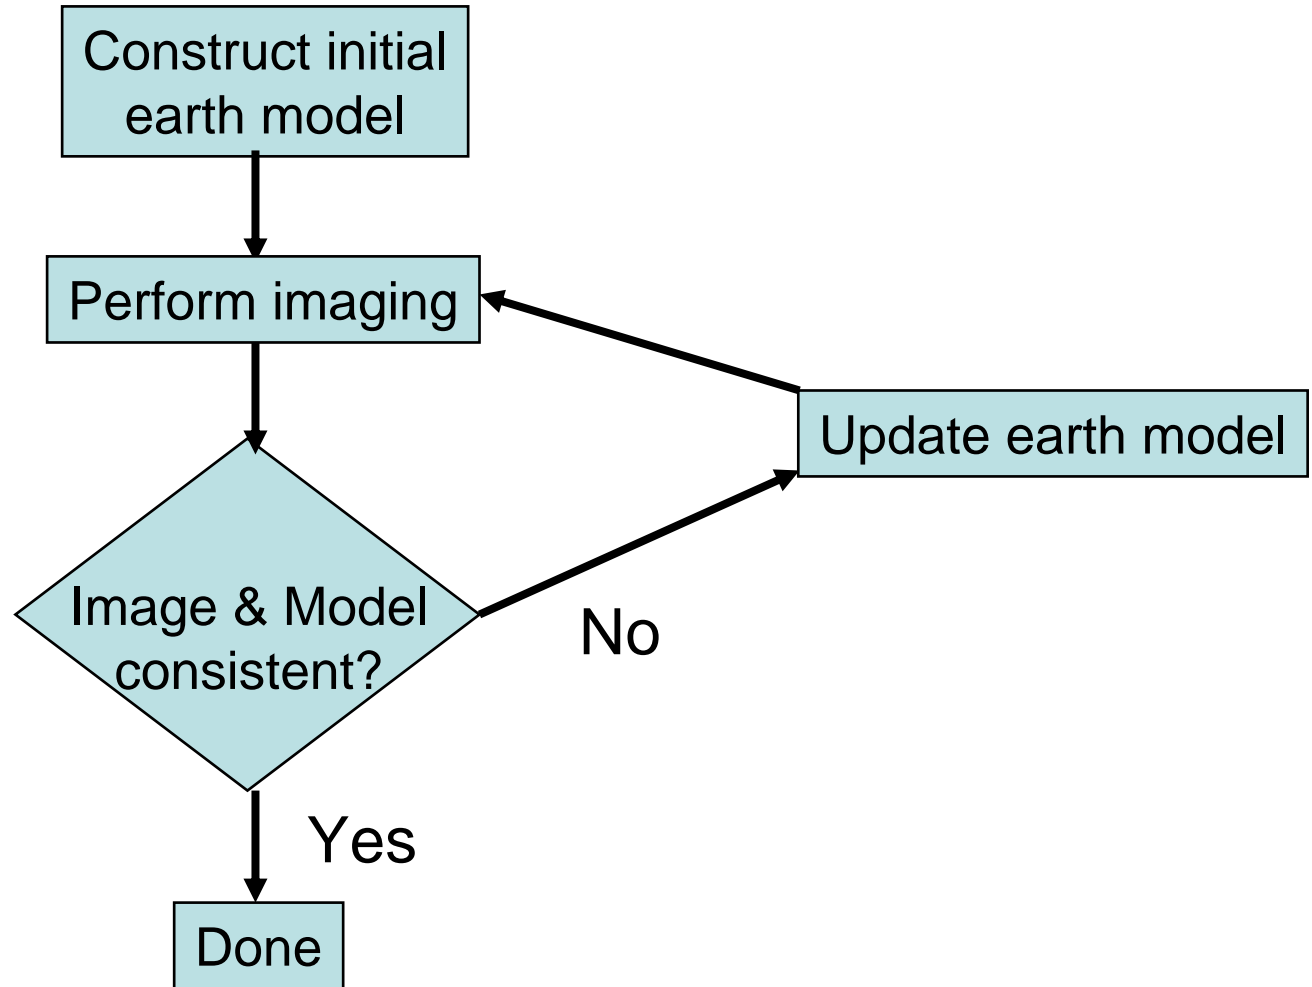


Seismic Imaging

An iterative process



Computation scales
as size of data and
image/model



- **Price-to-performance ratio improvement**
 - Want 10X to change platforms
 - Payback must more than cover effort & risk
 - Got 10X ten years ago in switching from supercomputers to PC clusters





- **Price-to-performance ratio improvement**
 - Want 10X to change platforms
 - Payback must more than cover effort & risk
 - Got 10X ten years ago in switching from supercomputers to PC clusters
 - Several years ago there were indicators we can get 10X or more on GPUs
 - Peak performance
 - Benchmarks
 - Simple prototype kernels



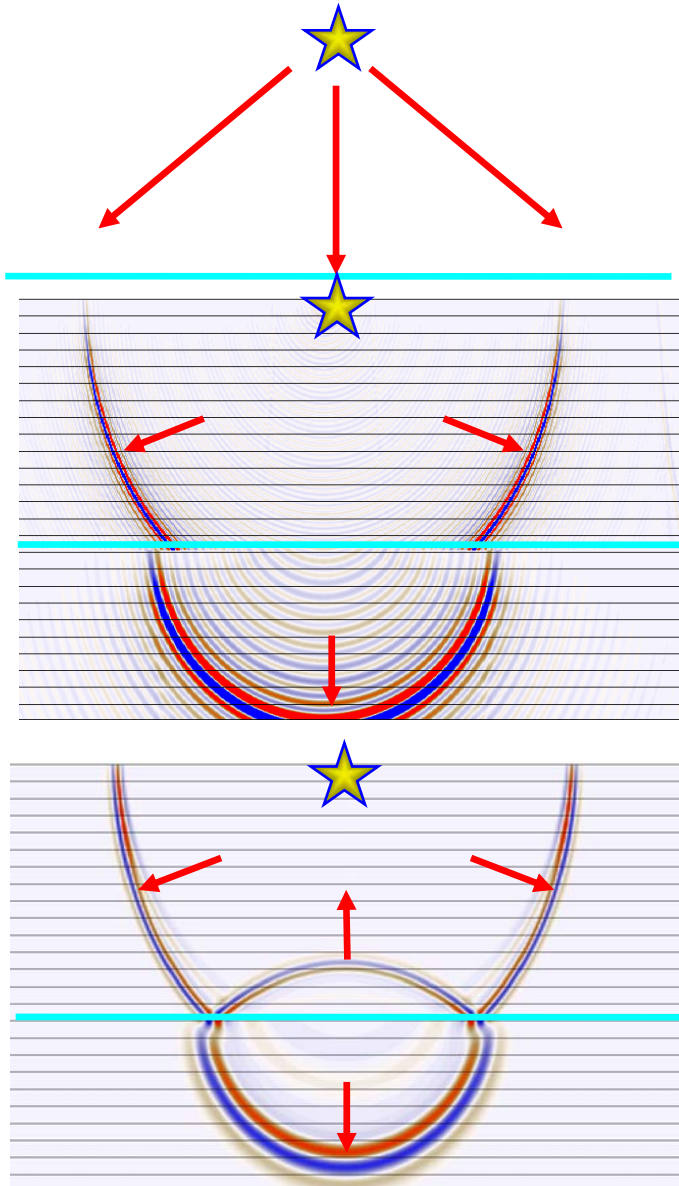
- **Ease of programming**

- Must be able to port, maintain & modify production codes (relatively) easily
 - These costs must be included
- Have tried Cg, Brook and Peakstream
 - All lacking in some aspect
- CUDA programming model straightforward
 - SIMD-like thread-based parallelism
 - In 1.5 days
 - Took "intro to CUDA" class
 - Wrote a working 2-D seismic modeling code
 - Programming memory hierarchy for optimization is the biggest challenge



- **Design GPU algorithm**
 - Optimize for memory hierarchy
 - Keep main data structures in GPU memory
- **Create prototype GPU kernel**
 - Include main computational characteristics
 - Test performance against CPU kernel
 - Iteratively refine prototype
- **Port full kernel & compare with CPU kernel**
 - Verify numerical results
 - Compare performance results
- **Incorporate into production code & system**

↙ Increasing computational cost ↘



- **Kirchhoff imaging**
 - High-frequency propagation
 - Ray or eikonal travel-times
- **"Wave-equation" imaging**
 - One-way propagation: $z \sim t$
 - Frequency-domain method
 - ADI (alternating direction implicit) finite difference
- **"Reverse-time" imaging**
 - Two-way propagation
 - Time-domain
 - Explicit finite-difference

- Based on the Kirchhoff integral

- Pre-compute coarse travel-times for propagation from surface locations to image points: $T(\vec{s}, \vec{x})$

- 4-D surface integral through a 5-D data set

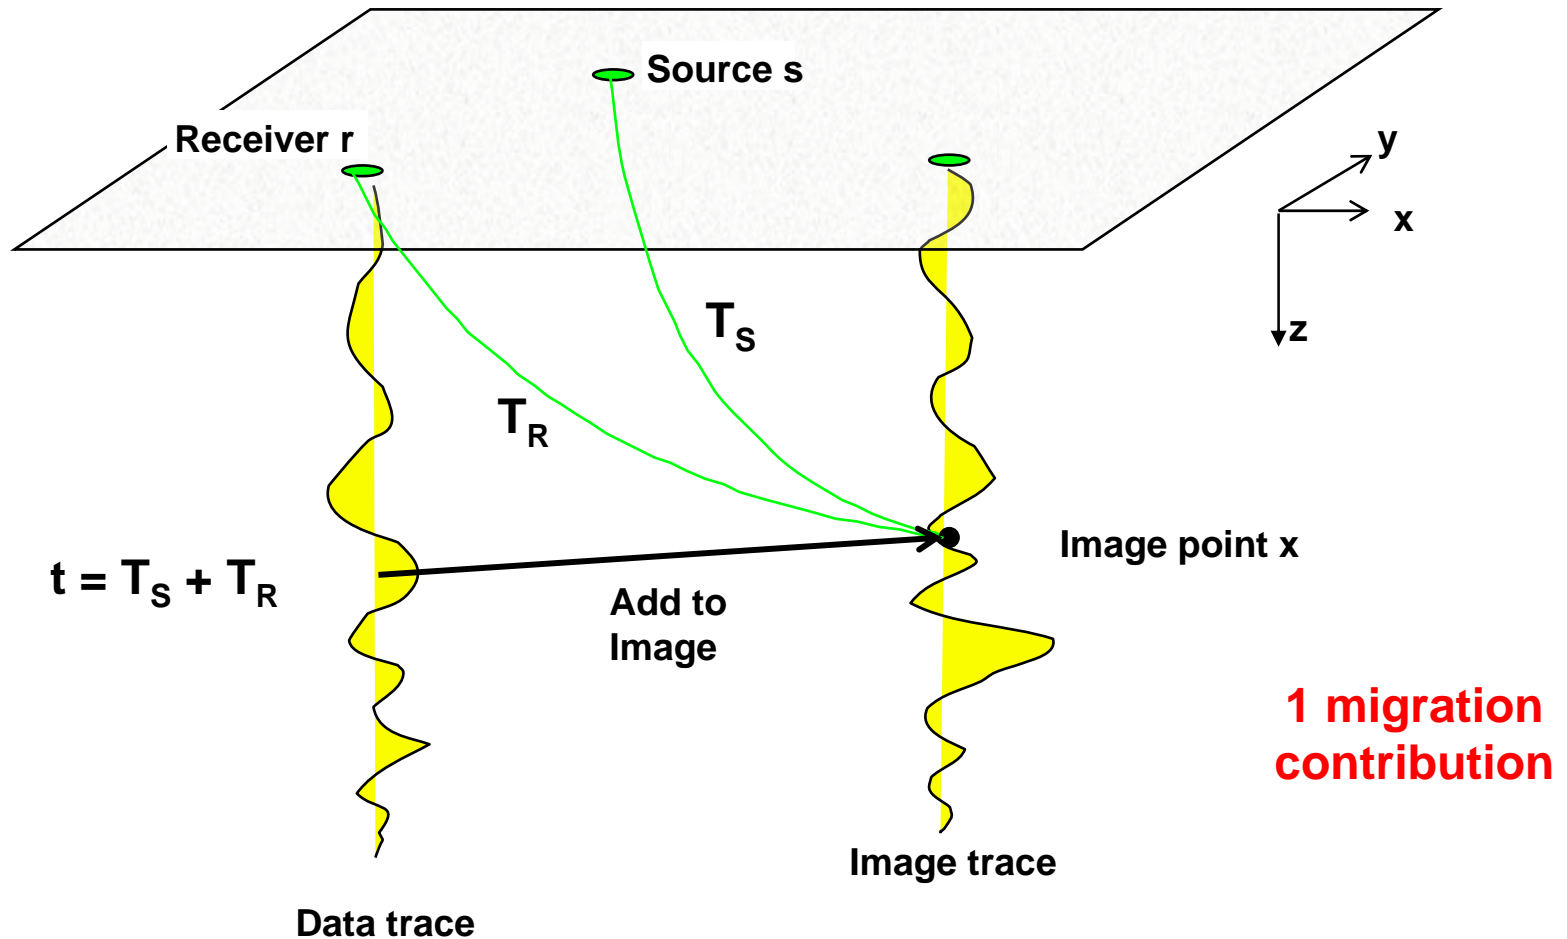
$$I(\vec{x}) = \iint d^2\mathbf{s} \, d^2\mathbf{r} \, D(\vec{s}, \vec{r}, t = T(\vec{s}, \vec{x}) + T(\vec{x}, \vec{r}))$$

- Computational complexity:

- $N_I \sim 10^9$ is the number of output image points
- $N_D \sim 10^8$ is the number of input data traces
- $f \sim 10$ is the number of cycles/point/trace
- $f N_I N_D \sim 10^{18}$ cycles ~ 10 CPU-years

Kirchhoff Imaging

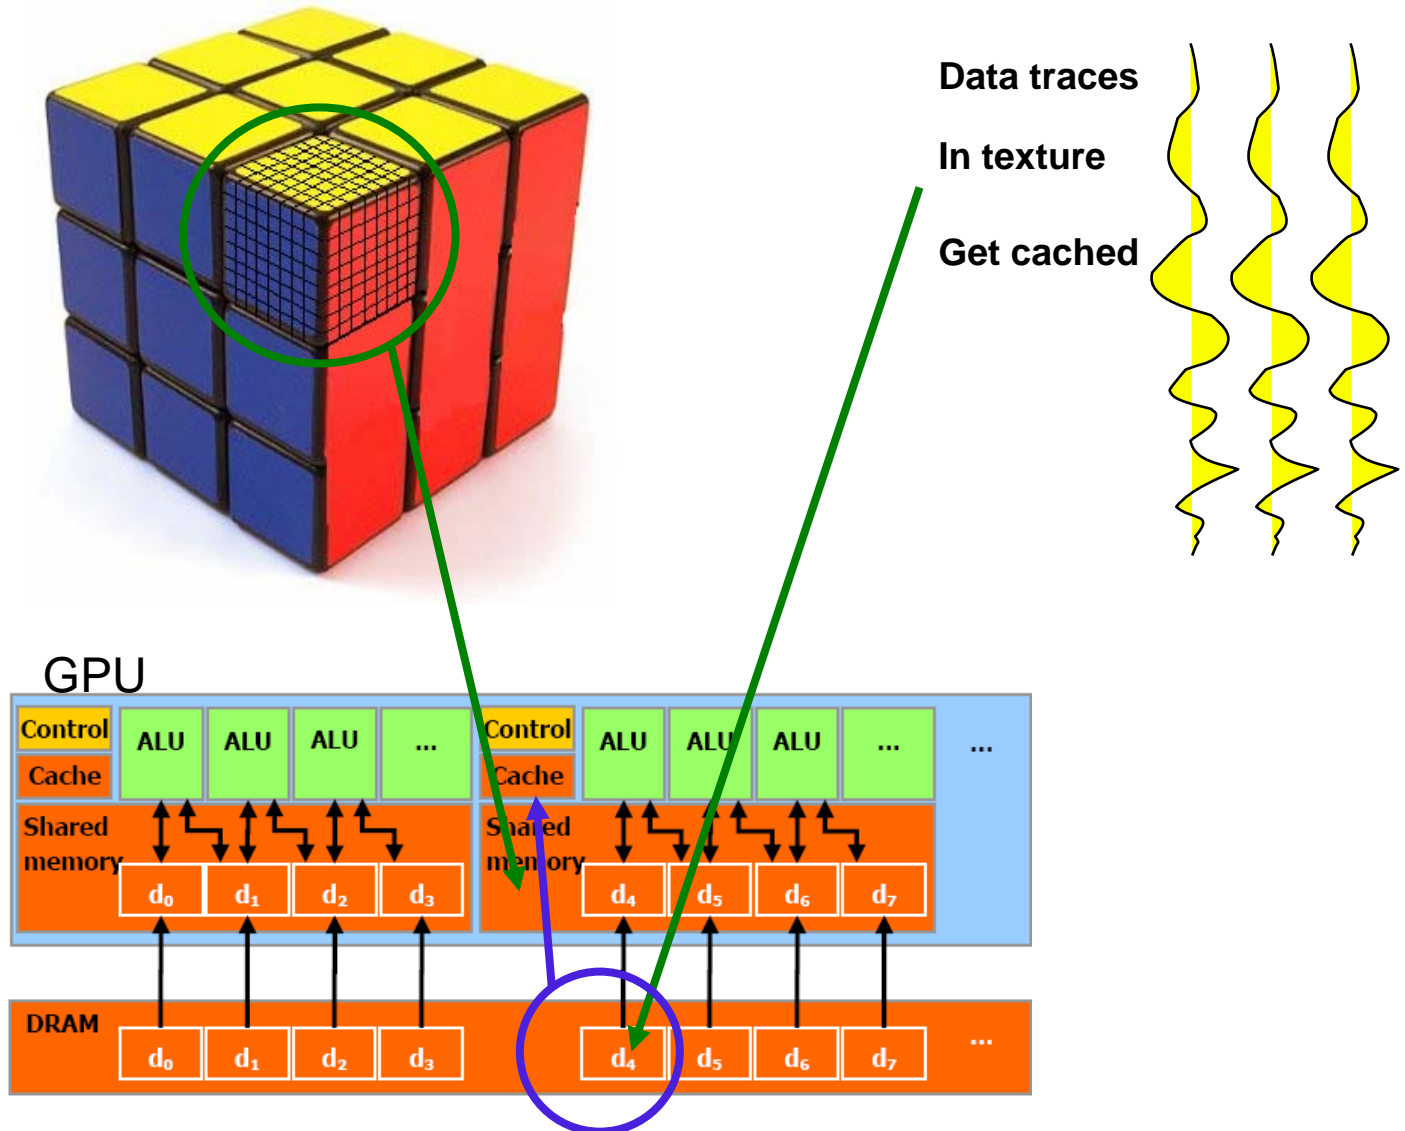
Computational kernel



$$I(\vec{\mathbf{x}}) = \sum_{\vec{\mathbf{s}}, \vec{\mathbf{r}}} D(\vec{\mathbf{s}}, \vec{\mathbf{r}}, t = T(\vec{\mathbf{s}}, \vec{\mathbf{x}}) + T(\vec{\mathbf{x}}, \vec{\mathbf{r}}))$$

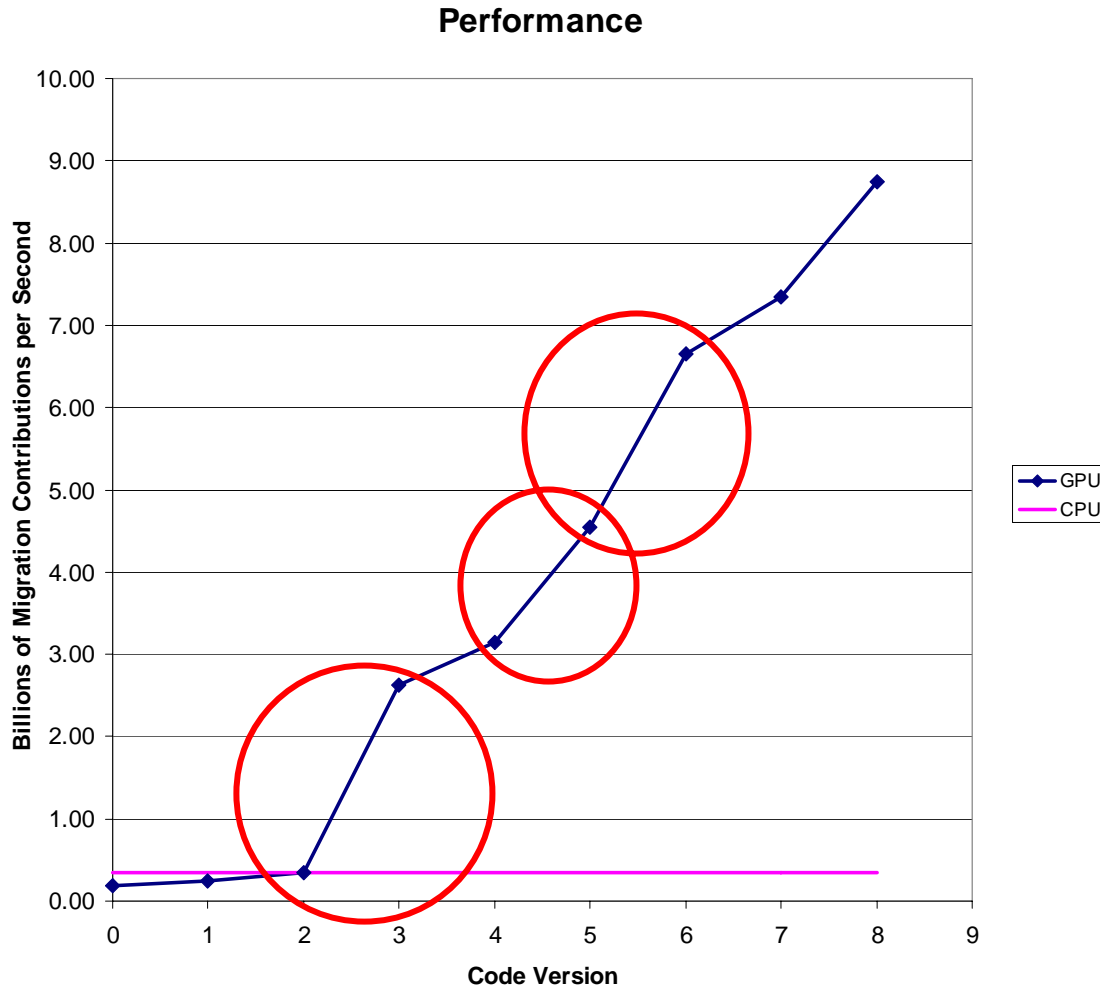
Kirchhoff Imaging

CUDA kernel



Kirchhoff Imaging

Kernel optimization



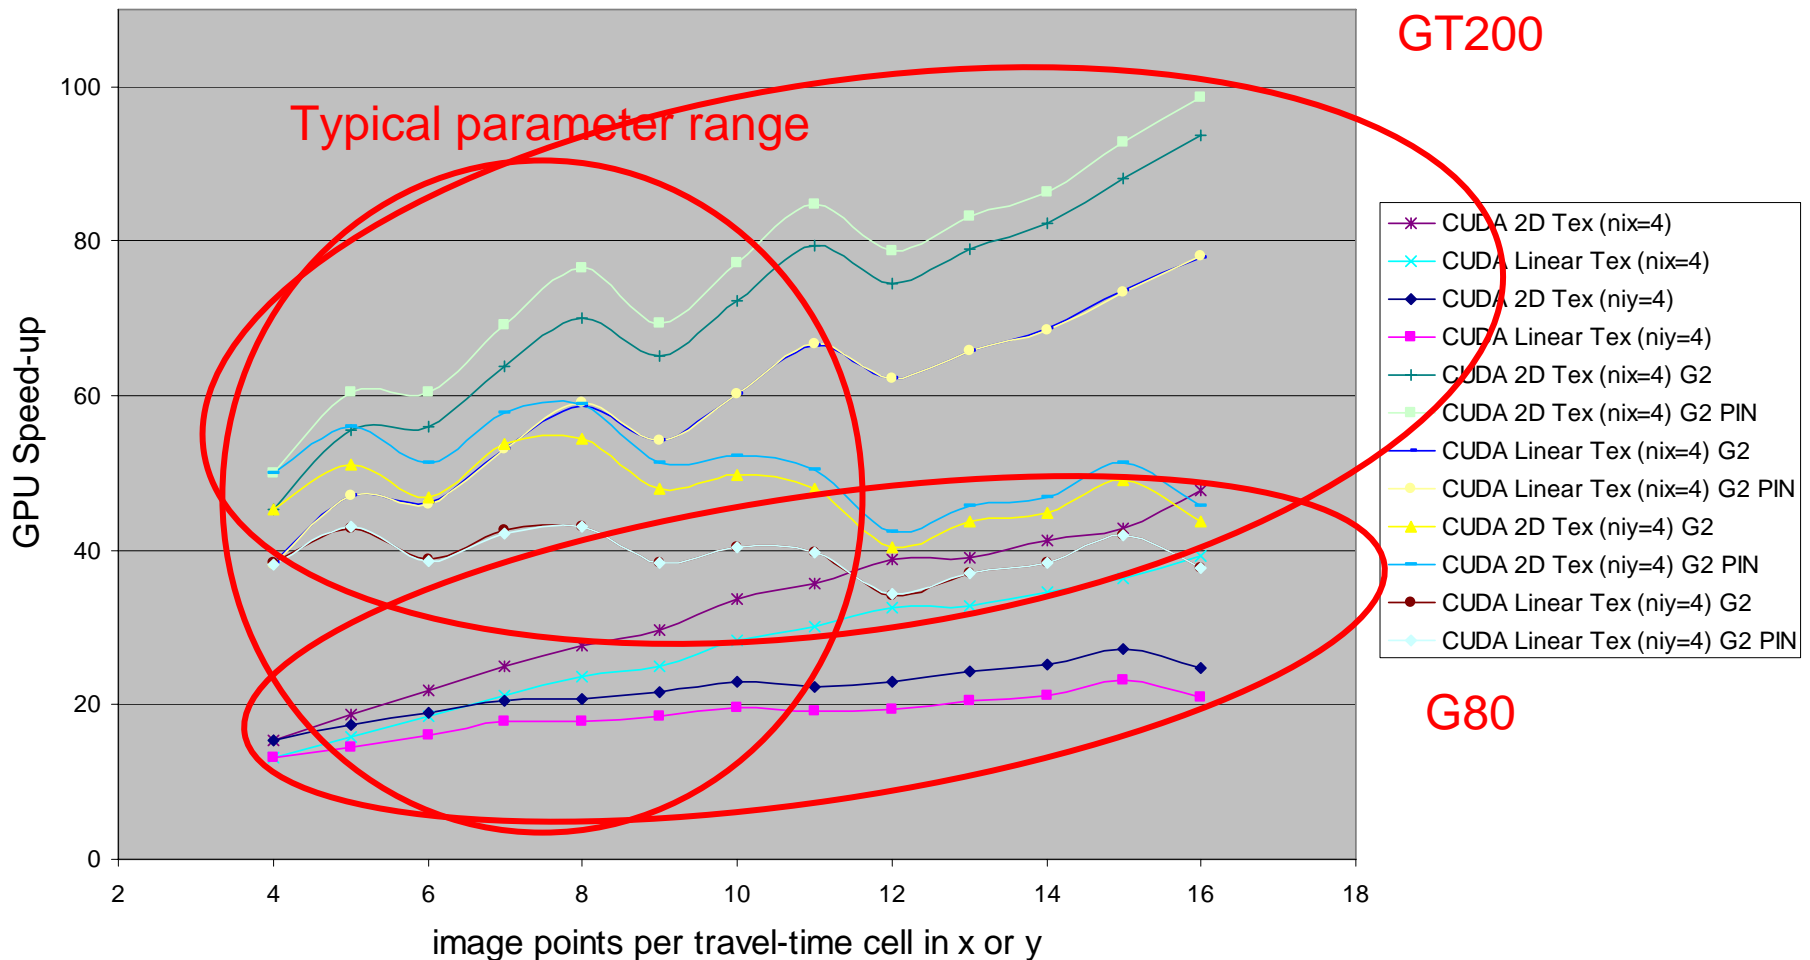
- 0 – Initial Kernel
- 1 – Used Texture Memory
- 2 – Used Shared Memory
- 3 – Global Memory Coalescing
- 4 – Decreased Data Trace Shared Memory Use
- 5 – Optimized Use of Shared Memory
- 6 – Consolidated “if” Statements, Eliminated or Substituted Some Math Operations
- 7 – Removed an “if” and “for”
- 8 – Used Texture Memory for Data-Trace Fetch

Kirchhoff Imaging

Kernel performance



GPU-to-CPU Performance Ratio

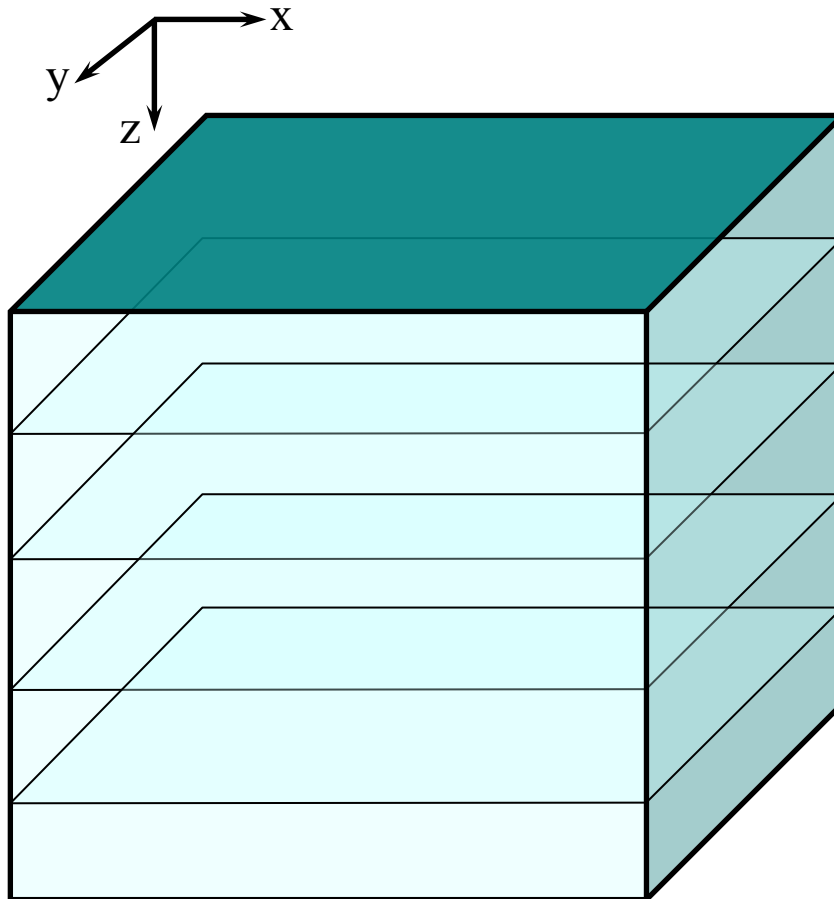


- GPU kernel incorporated into production code
 - Large kernel speed-ups results in “CPU overhead” for task setup dominating GPU production runs
- Further optimizations
 - create GPU kernels for most “overhead” components
 - optimized left-over CPU code (which helps CPU version also)

Time (hr)	Set-up	Kernel	Total	Speed-up
Original CPU code	5	20	25	
Main GPU kernel	5	0.5	5.5	5
Further optimizations	0.5	0.5	1	25

"Wave-equation" Imaging

One-way propagation



- Based on scalar wave equation

$$\frac{1}{V(\vec{\mathbf{x}})^2} \frac{\partial^2 P}{\partial t^2} = \nabla^2 P$$

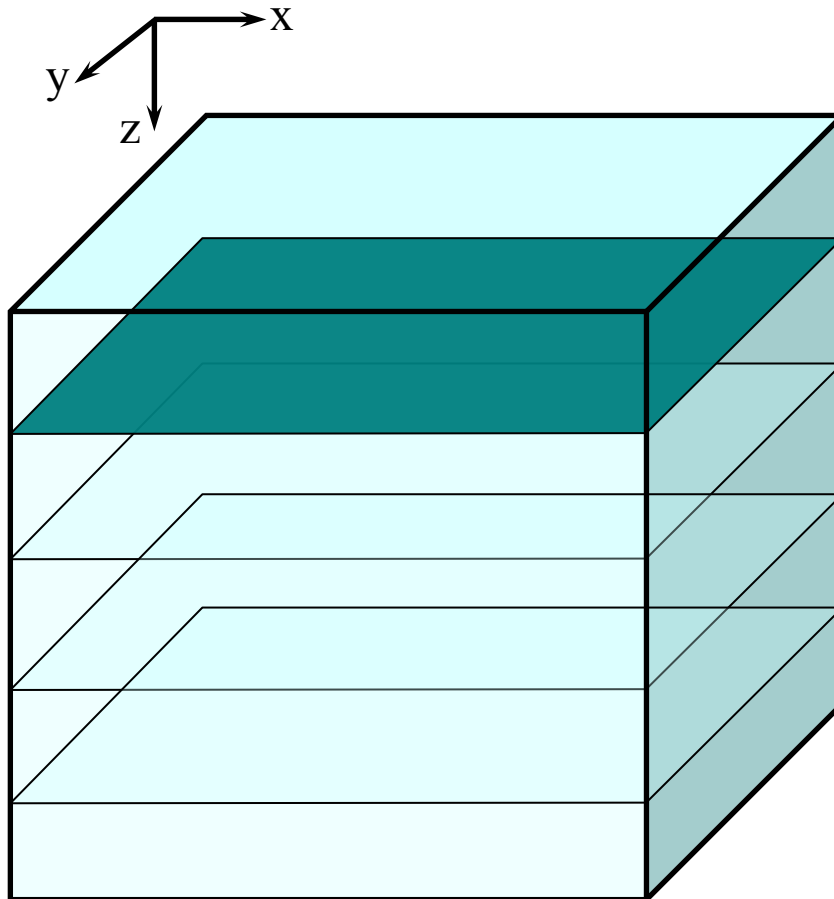
- Frequency-domain
- Preferred direction of propagation: $z \sim t$

$$\frac{\partial P}{\partial z} = \frac{\pm i\omega}{V(\vec{\mathbf{x}})} \sqrt{1 + \frac{V^2(\vec{\mathbf{x}})}{\omega^2} \left(\frac{\partial^2}{\partial x^2} + \frac{\partial^2}{\partial y^2} \right)} P$$

- Evolution in depth

"Wave-equation" Imaging

One-way propagation



- Based on scalar wave equation

$$\frac{1}{V(\vec{\mathbf{x}})^2} \frac{\partial^2 P}{\partial t^2} = \nabla^2 P$$

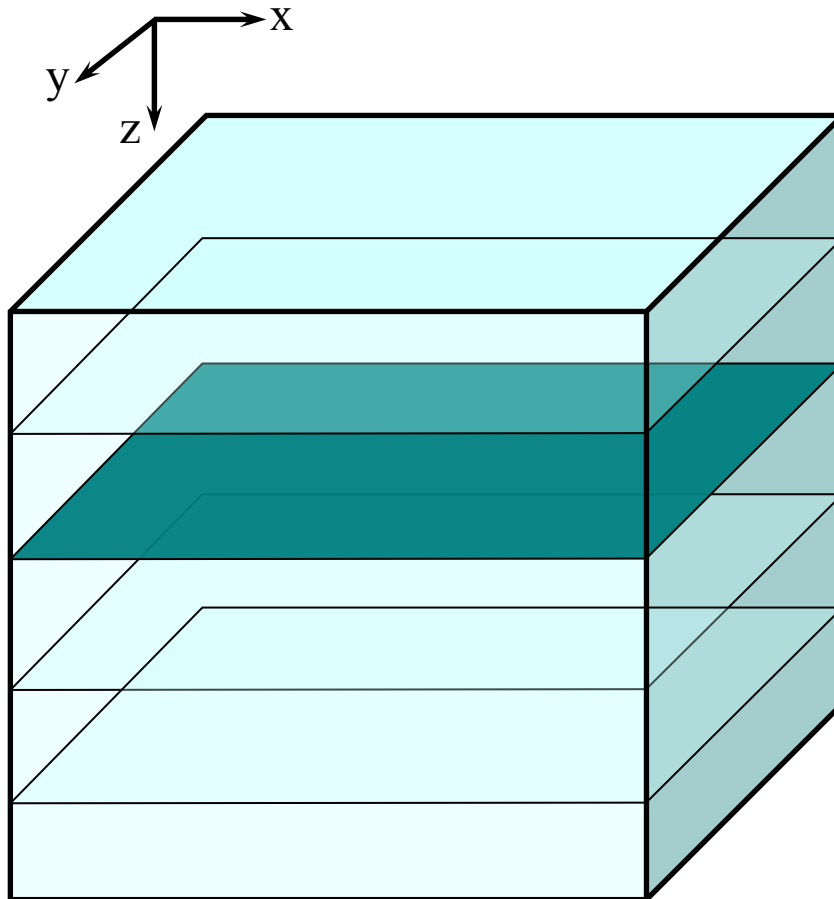
- Frequency-domain
- Preferred direction of propagation: $z \sim t$

$$\frac{\partial P}{\partial z} = \frac{\pm i\omega}{V(\vec{\mathbf{x}})} \sqrt{1 + \frac{V^2(\vec{\mathbf{x}})}{\omega^2} \left(\frac{\partial^2}{\partial x^2} + \frac{\partial^2}{\partial y^2} \right)} P$$

- Evolution in depth

"Wave-equation" Imaging

One-way propagation



- Based on scalar wave equation

$$\frac{1}{V(\vec{\mathbf{x}})^2} \frac{\partial^2 P}{\partial t^2} = \nabla^2 P$$

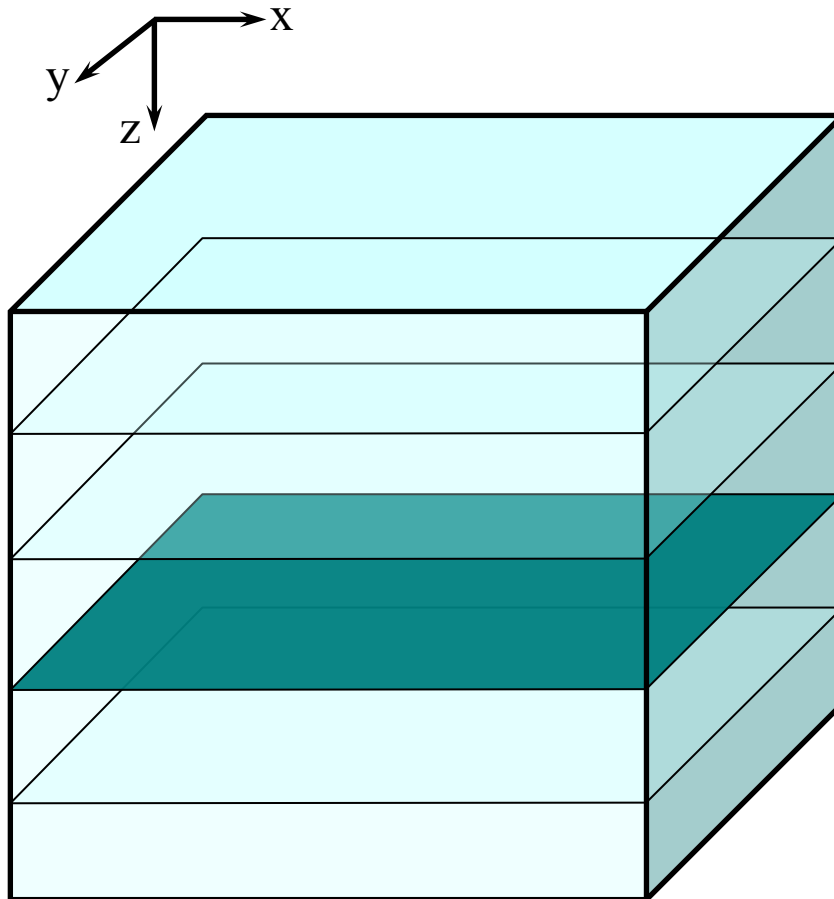
- Frequency-domain
- Preferred direction of propagation: $z \sim t$

$$\frac{\partial P}{\partial z} = \frac{\pm i\omega}{V(\vec{\mathbf{x}})} \sqrt{1 + \frac{V^2(\vec{\mathbf{x}})}{\omega^2} \left(\frac{\partial^2}{\partial x^2} + \frac{\partial^2}{\partial y^2} \right)} P$$

- Evolution in depth

"Wave-equation" Imaging

One-way propagation



- Based on scalar wave equation

$$\frac{1}{V(\vec{\mathbf{x}})^2} \frac{\partial^2 P}{\partial t^2} = \nabla^2 P$$

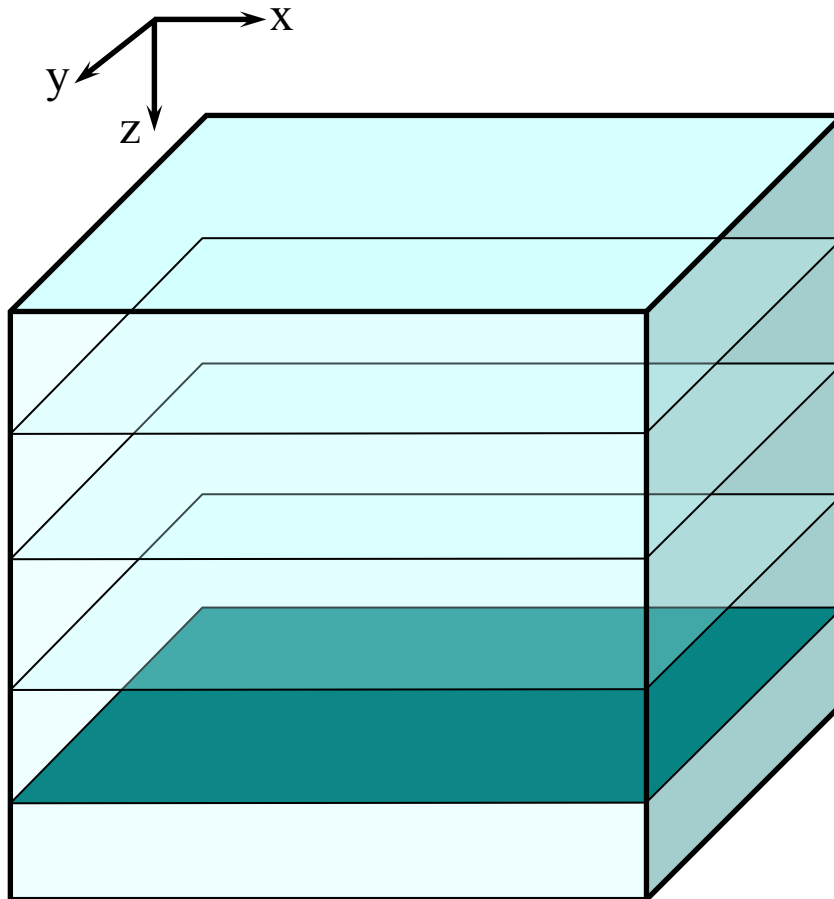
- Frequency-domain
- Preferred direction of propagation: $z \sim t$

$$\frac{\partial P}{\partial z} = \frac{\pm i\omega}{V(\vec{\mathbf{x}})} \sqrt{1 + \frac{V^2(\vec{\mathbf{x}})}{\omega^2} \left(\frac{\partial^2}{\partial x^2} + \frac{\partial^2}{\partial y^2} \right)} P$$

- Evolution in depth

"Wave-equation" Imaging

One-way propagation



- Based on scalar wave equation

$$\frac{1}{V(\vec{\mathbf{x}})^2} \frac{\partial^2 P}{\partial t^2} = \nabla^2 P$$

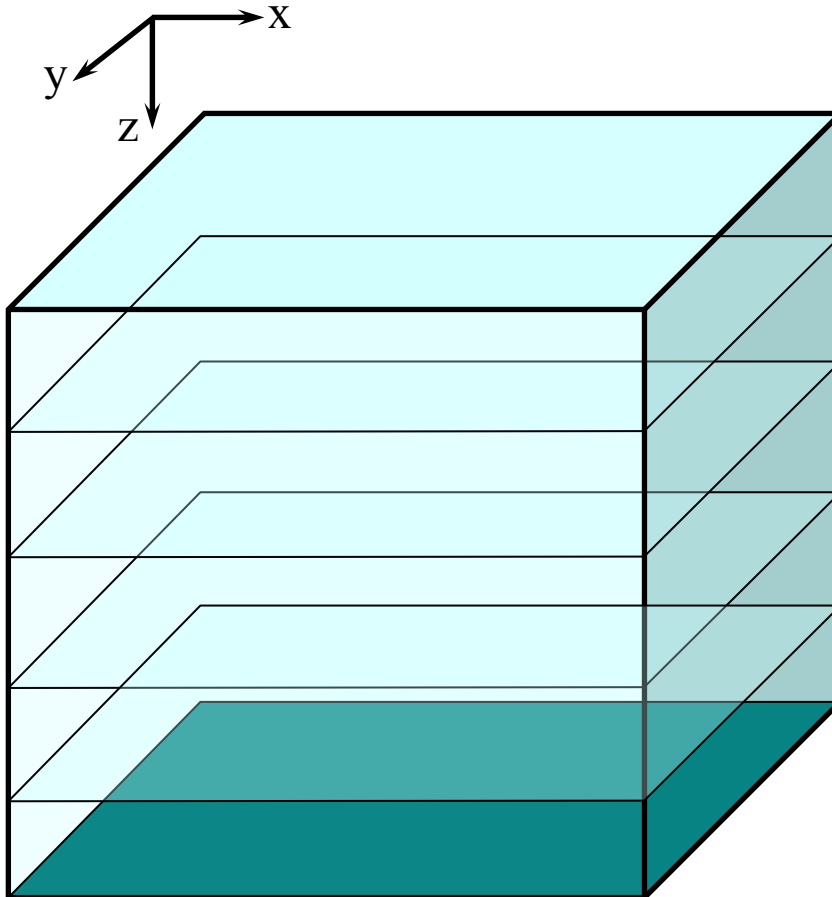
- Frequency-domain
- Preferred direction of propagation: $z \sim t$

$$\frac{\partial P}{\partial z} = \frac{\pm i\omega}{V(\vec{\mathbf{x}})} \sqrt{1 + \frac{V^2(\vec{\mathbf{x}})}{\omega^2} \left(\frac{\partial^2}{\partial x^2} + \frac{\partial^2}{\partial y^2} \right)} P$$

- Evolution in depth

"Wave-equation" Imaging

One-way propagation



- Based on scalar wave equation

$$\frac{1}{V(\vec{\mathbf{x}})^2} \frac{\partial^2 P}{\partial t^2} = \nabla^2 P$$

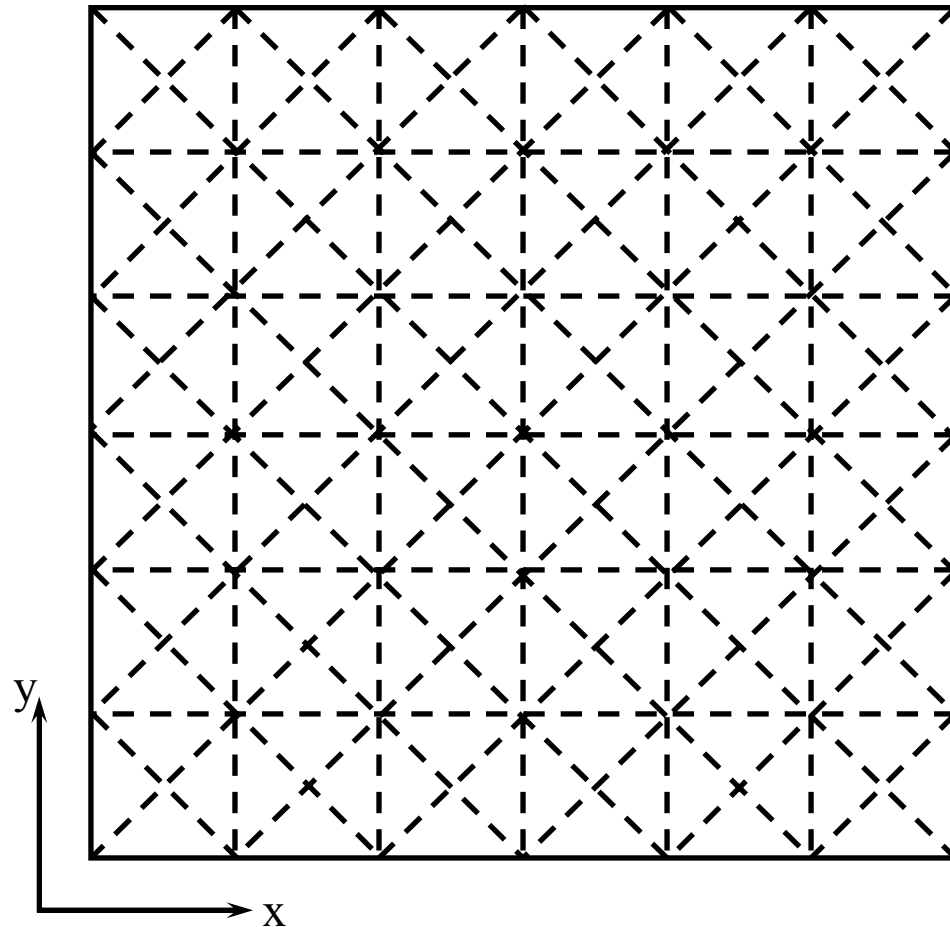
- Frequency-domain
- Preferred direction of propagation: $z \sim t$

$$\frac{\partial P}{\partial z} = \frac{\pm i\omega}{V(\vec{\mathbf{x}})} \sqrt{1 + \frac{V^2(\vec{\mathbf{x}})}{\omega^2} \left(\frac{\partial^2}{\partial x^2} + \frac{\partial^2}{\partial y^2} \right)} P$$

- Evolution in depth

“Wave-equation” Imaging

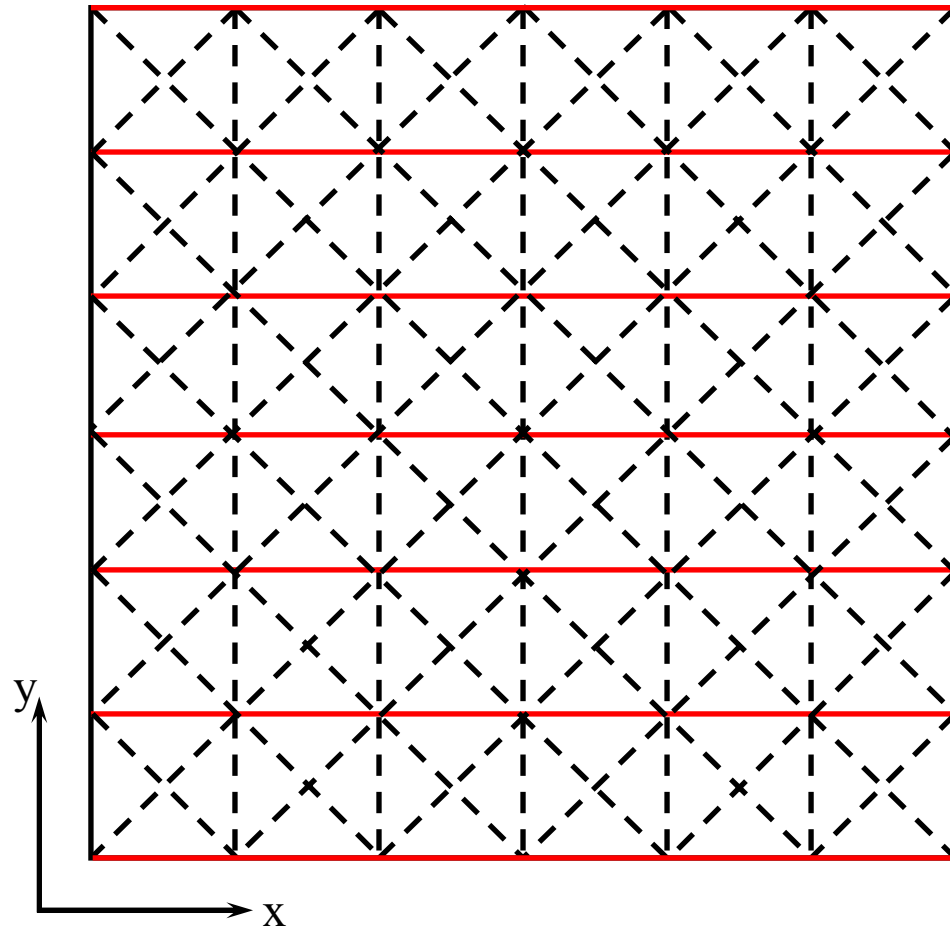
One-way propagation



- Evolution eqn uses
 - Continued fractions
 - Operator splitting
 - ADI finite difference
- Each depth step requires applying four operators
 - Along x
 - Along y
 - Along $x + y$
 - Along $x - y$

“Wave-equation” Imaging

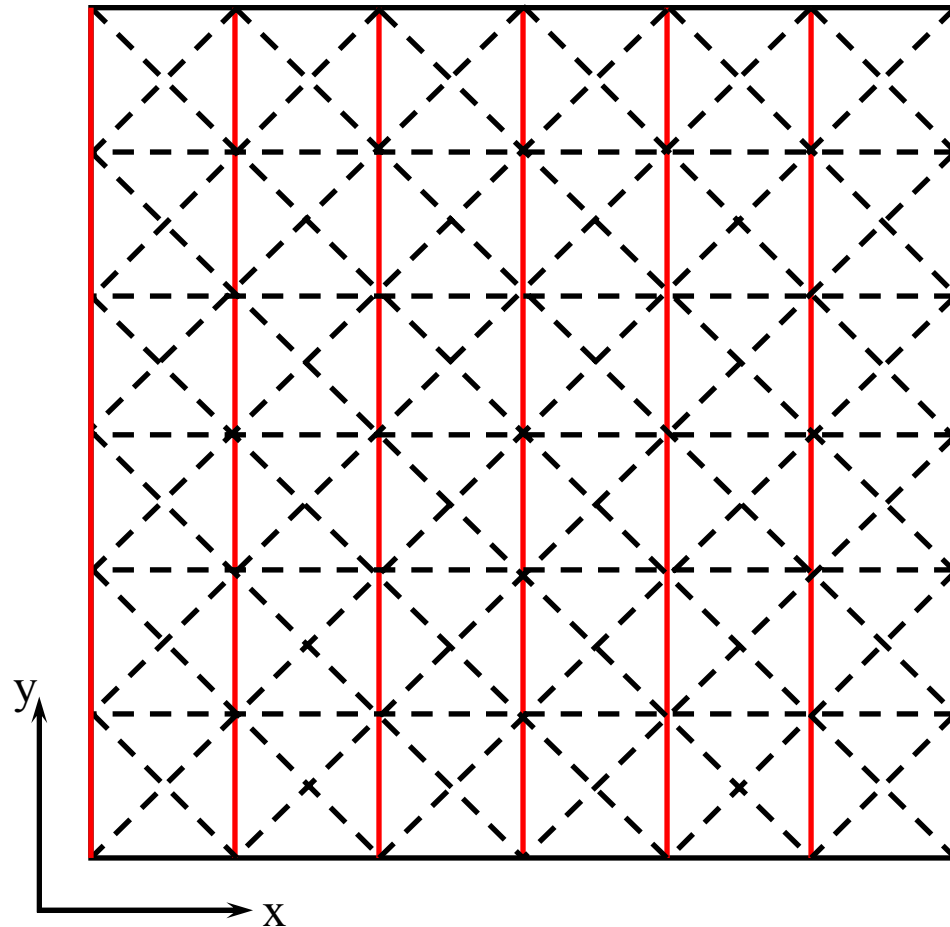
One-way propagation



- Evolution eqn uses
 - Continued fractions
 - Operator splitting
 - ADI finite difference
- Each depth step requires applying four operators
 - Along x
 - Along y
 - Along $x + y$
 - Along $x - y$

“Wave-equation” Imaging

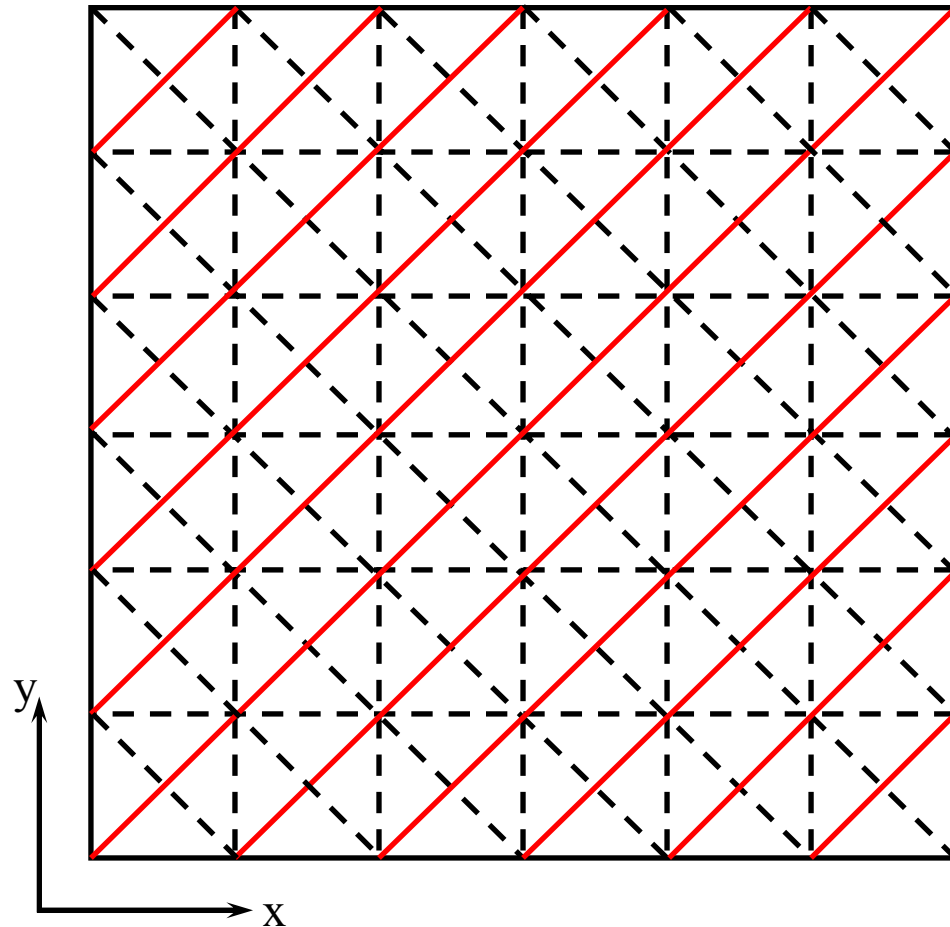
One-way propagation



- Evolution eqn uses
 - Continued fractions
 - Operator splitting
 - ADI finite difference
- Each depth step requires applying four operators
 - Along x
 - Along y
 - Along $x + y$
 - Along $x - y$

“Wave-equation” Imaging

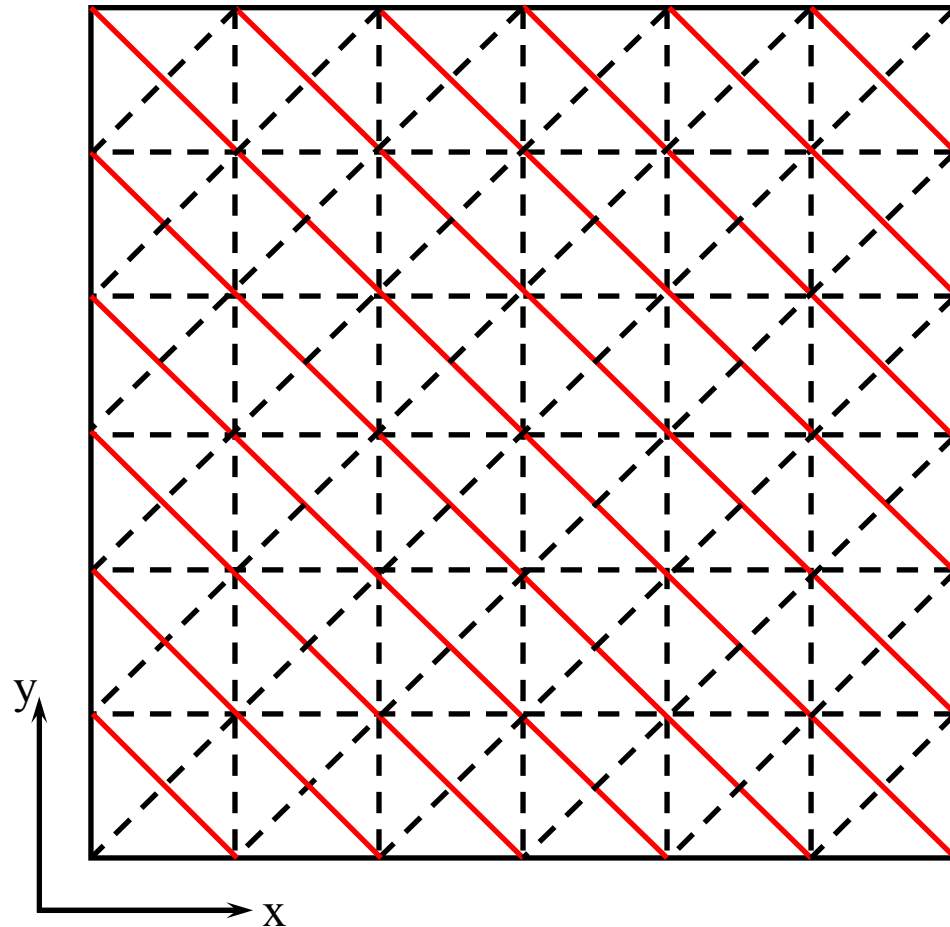
One-way propagation



- Evolution eqn uses
 - Continued fractions
 - Operator splitting
 - ADI finite difference
- Each depth step requires applying four operators
 - Along x
 - Along y
 - Along $x + y$
 - Along $x - y$

“Wave-equation” Imaging

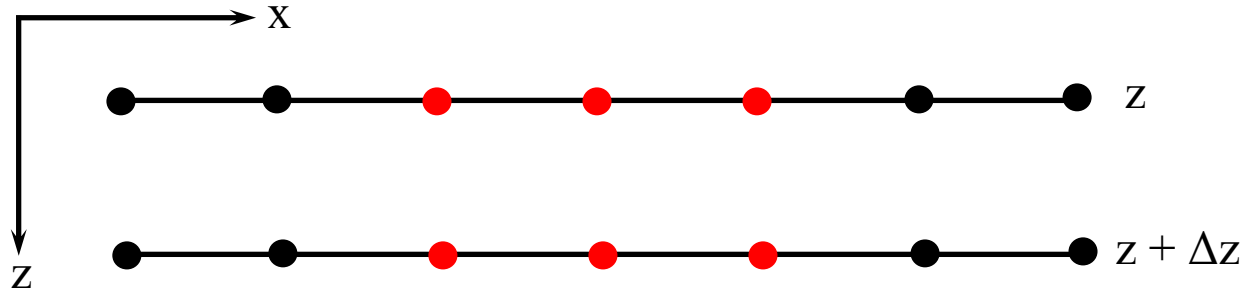
One-way propagation



- Evolution eqn uses
 - Continued fractions
 - Operator splitting
 - ADI finite difference
- Each depth step requires applying four operators
 - Along x
 - Along y
 - Along $x + y$
 - Along $x - y$

"Wave-equation" Imaging

Implicit complex tri-diagonal linear systems

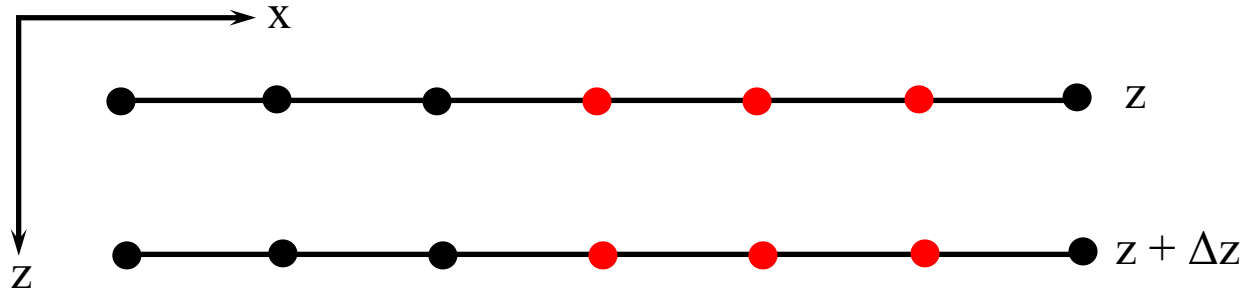


$$AP_{x-\Delta x, y}^{z+\Delta z} + (1-2A)P_{x, y}^{z+\Delta z} + AP_{x+\Delta x, y}^{z+\Delta z} =$$

$$A^*P_{x-\Delta x, y}^z + (1-2A^*)P_{x, y}^z + A^*P_{x+\Delta x, y}^z$$

“Wave-equation” Imaging

Implicit complex tri-diagonal linear systems



$$AP_{x-\Delta x, y}^{z+\Delta z} + (1-2A)P_{x, y}^{z+\Delta z} + AP_{x+\Delta x, y}^{z+\Delta z} =$$

$$A^*P_{x-\Delta x, y}^z + (1-2A^*)P_{x, y}^z + A^*P_{x+\Delta x, y}^z$$

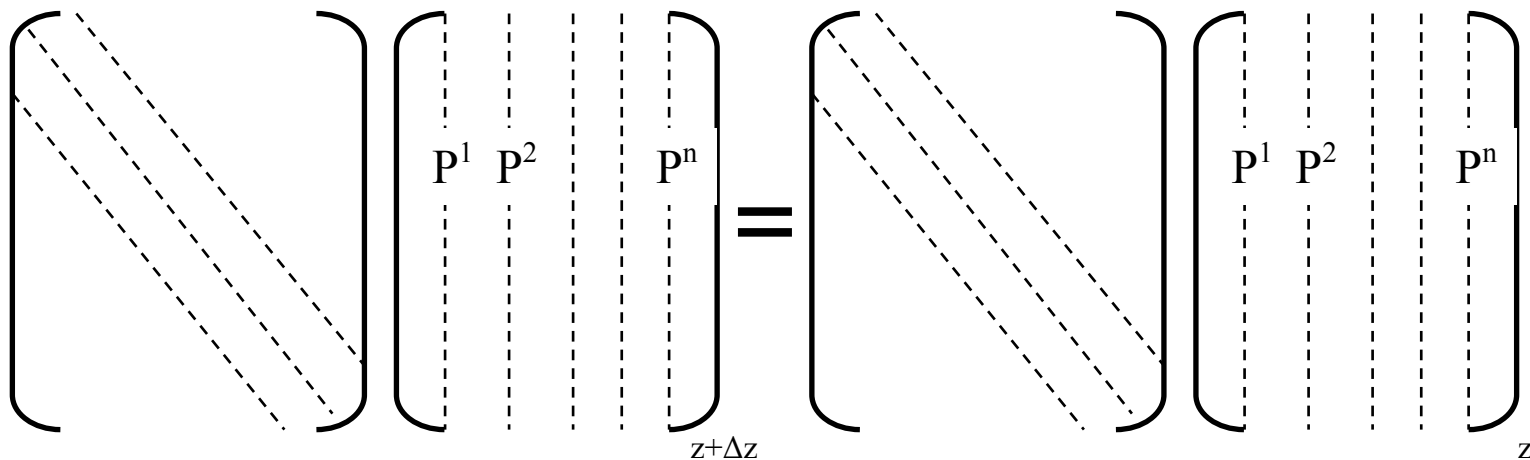
The evaluation and solution of these complex tri-diagonal systems dominates the computational cost of our wave-equation imaging code.

“Wave-equation” Imaging

Low level parallelism



- Common work between shot-records
 - Calculating the coefficients of the matrices
 - Dependent on frequency & local velocity
 - Part of the solving of the tri-diagonal system
- Parallelize over shot-records in the kernel



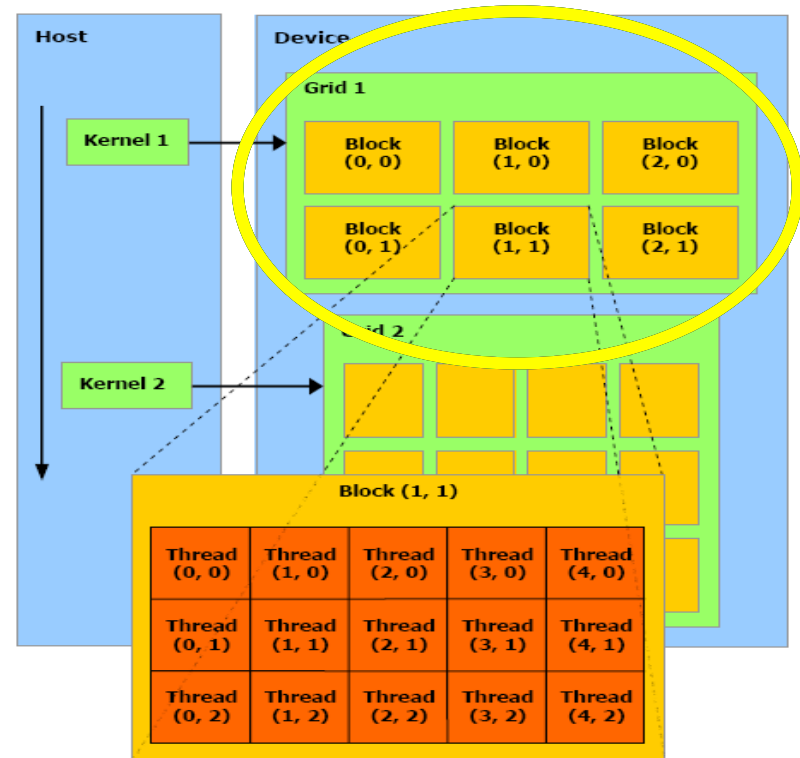
“Wave-equation” Imaging

CUDA kernels



- Separate kernels for each operator

- x y $x+y$ and $x-y$

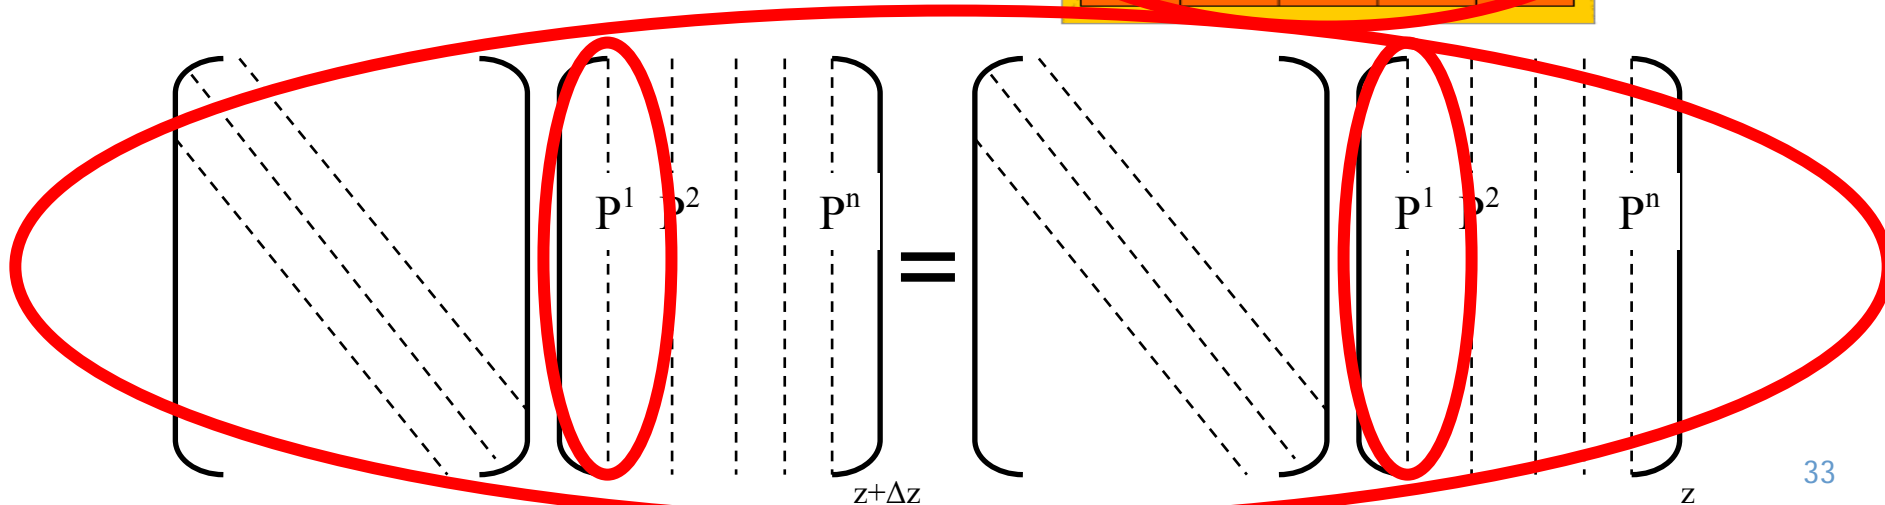
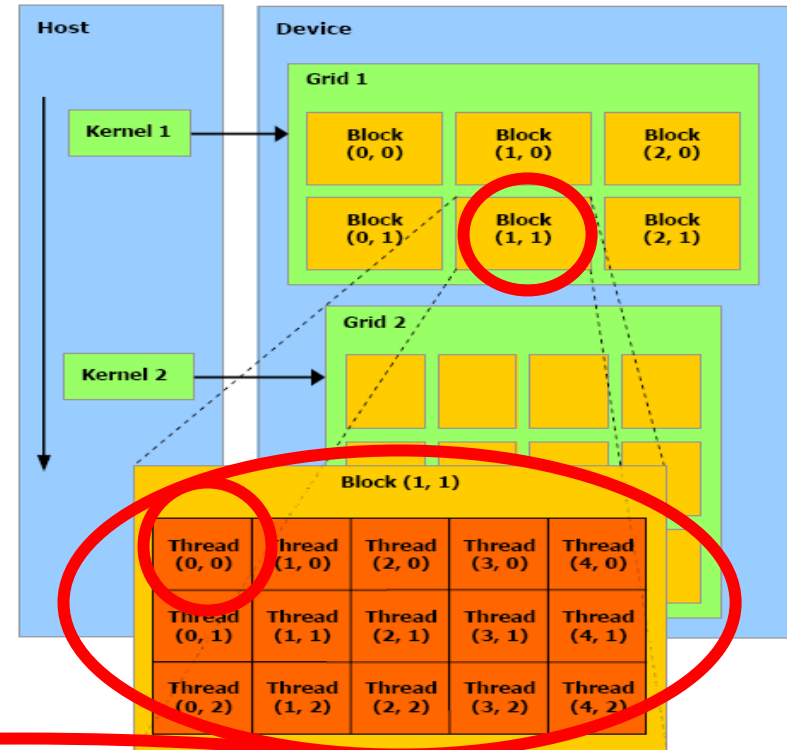


“Wave-equation” Imaging

CUDA kernels

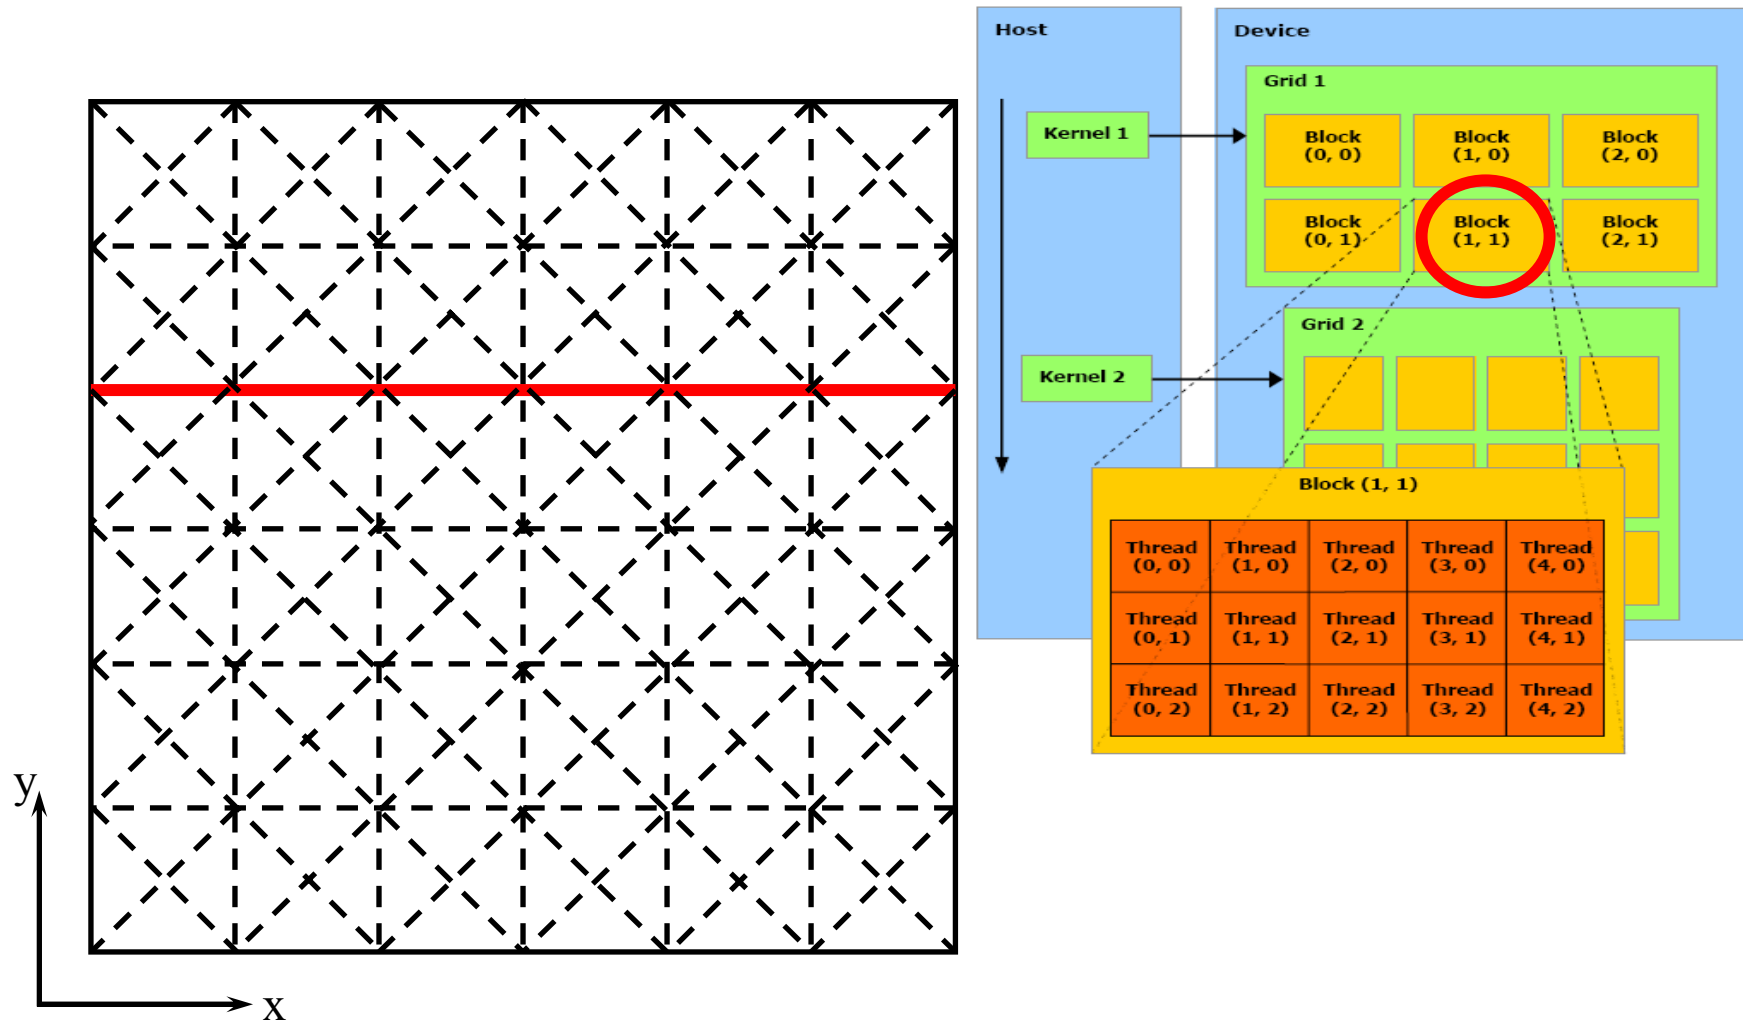


- Separate kernels for each operator
 - x , y , $x+y$ and $x-y$



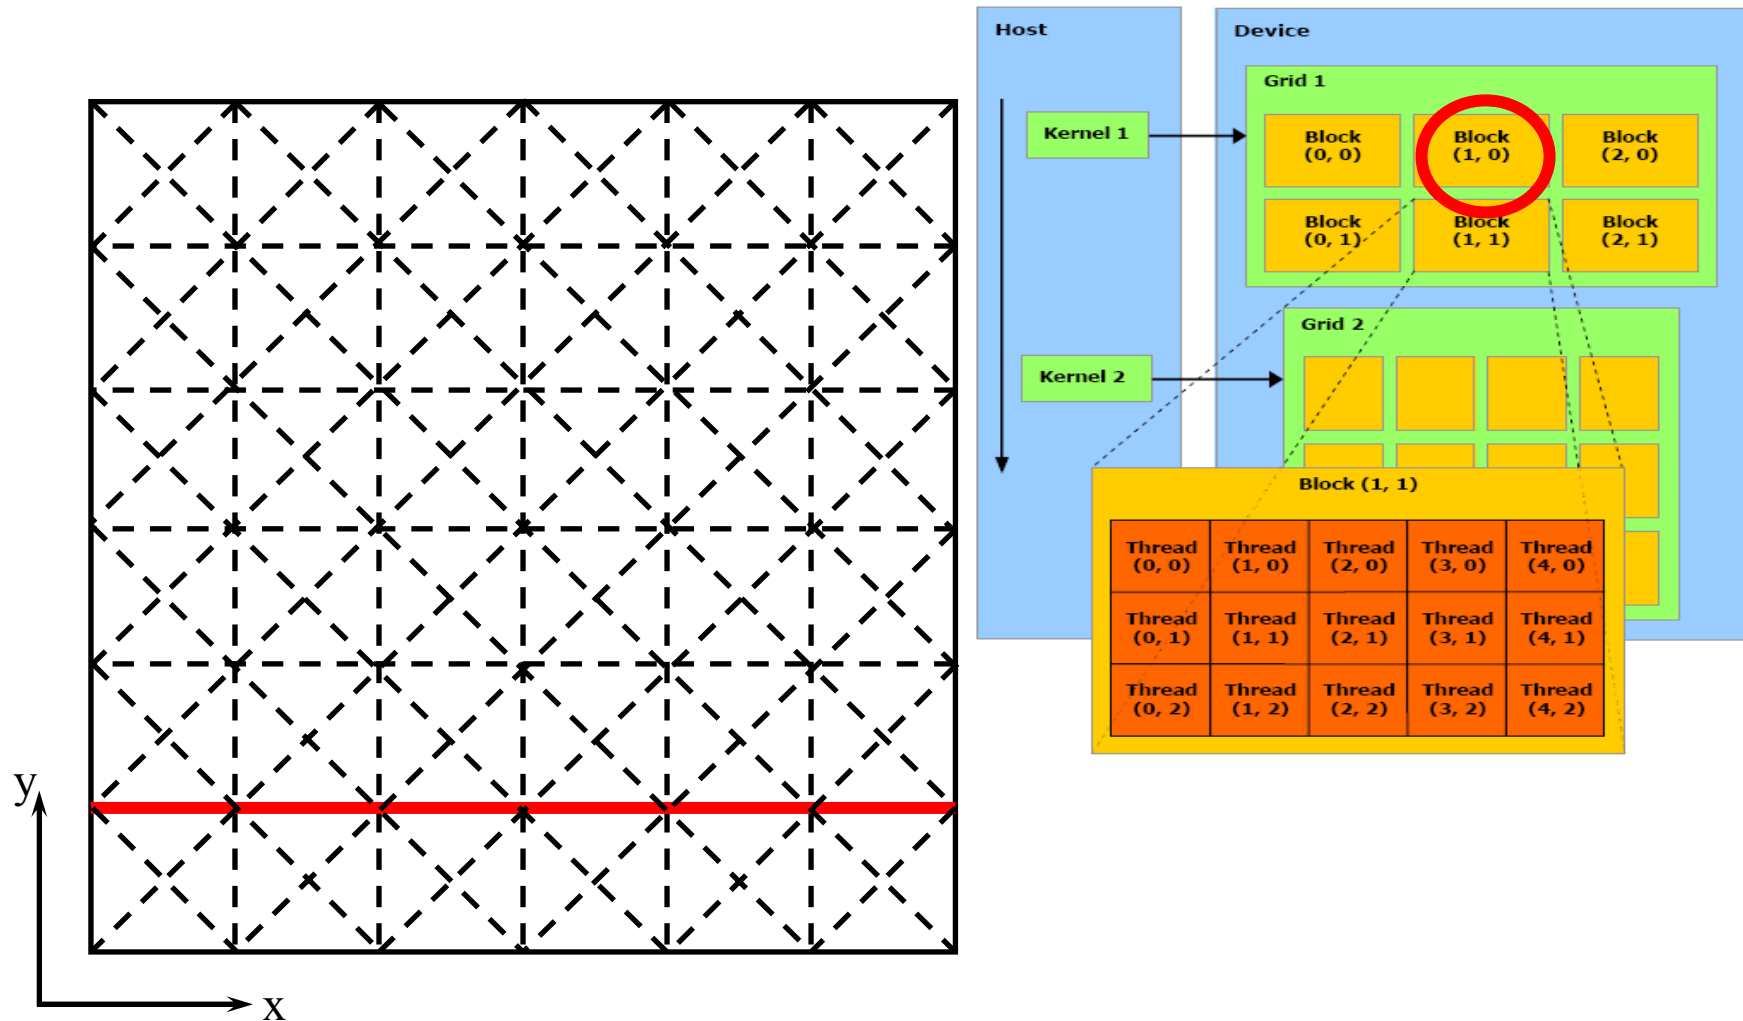
“Wave-equation” Imaging

CUDA kernels



“Wave-equation” Imaging

CUDA kernels



“Wave-equation” Imaging

Performance



- **Production CPU kernel**
 - Performance: 15 – 50 Mpoints/sec
- **Prototype CUDA kernel**
 - Single tri-diagonal system
 - Constant coefficients
 - Performance: 700 Mpoints/sec
- **Production CUDA kernels**
 - Single kernel handles x , y , $x+y$ & $x-y$ operators
 - Several kernels calculate coefficients
 - Performance: 300-500 Mpoints/sec

“Reverse-time” Imaging

Two-way propagation

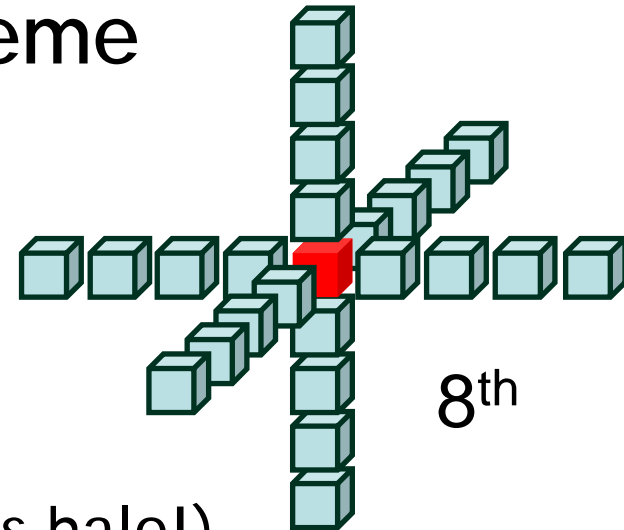


- Based on the scalar wave equation

$$\frac{1}{V(\mathbf{x})^2} \frac{\partial^2 P}{\partial t^2} = \nabla^2 P$$

- Explicit finite-difference scheme

- 2nd order in time
- Variable order in space: 6th - 16th
 - Most of the computation
- Bandwidth
 - Read $P(x, t)$, $P(x, t-dt)$ & $V(x)$ (plus halo!)
 - Write $P(x, t+dt)$
 - Max performance is 4+ Gpt/s

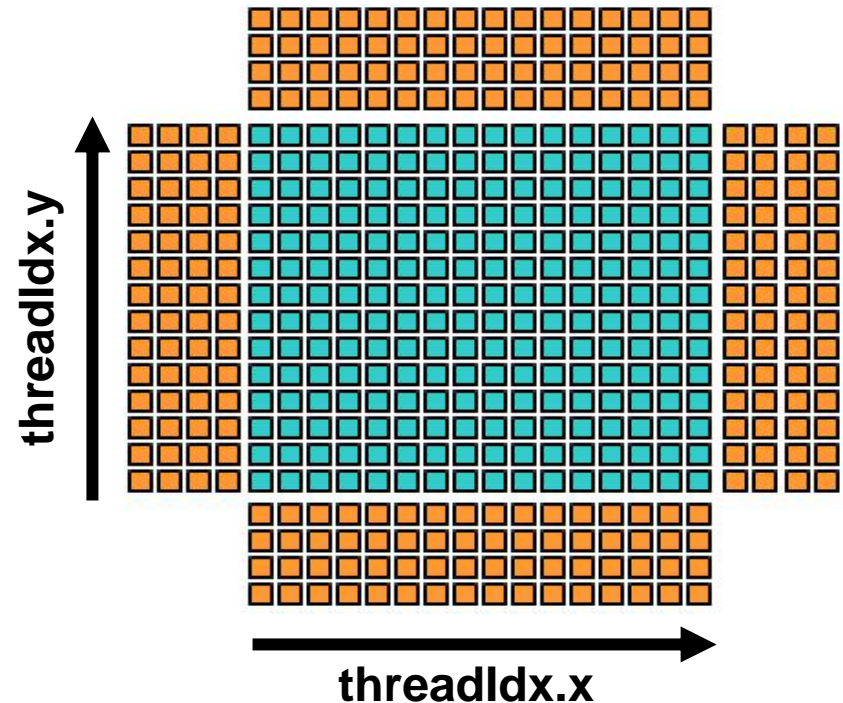


“Reverse-time” Imaging

CUDA algorithm



- Paulius's algorithm
 - Each thread specifies an (x,y) point, marching in z .
 - Each thread block handles a 2-D rectangle.
 - Each 2-D slice + halo is read into shared memory.
 - Threads in a block re-use these values.

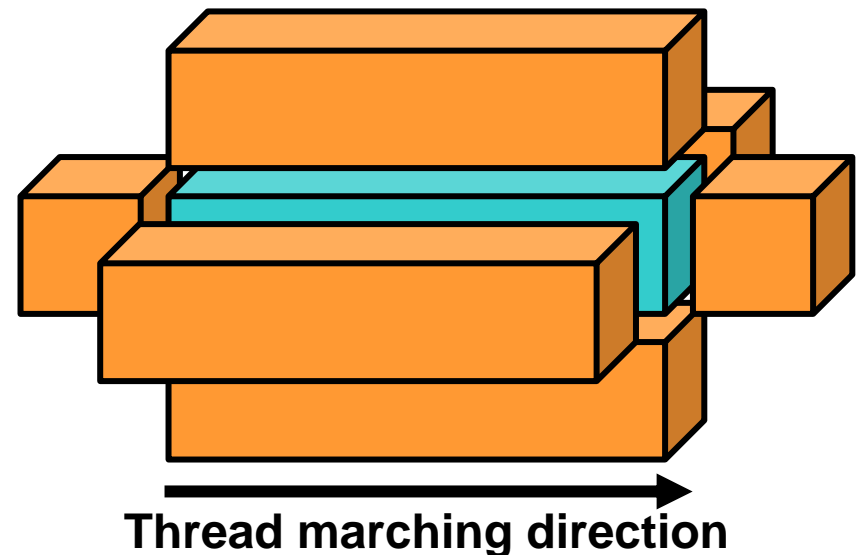


“Reverse-time” Imaging

CUDA algorithm



- Paulius’s algorithm
 - Threads in a block march in z , storing values “in-front” & “behind” in registers.
 - Number of registers limits the block size and the core-to-halo ratio.
 - GPU performance is predictable: 2.5 - 3 Gpt/s.

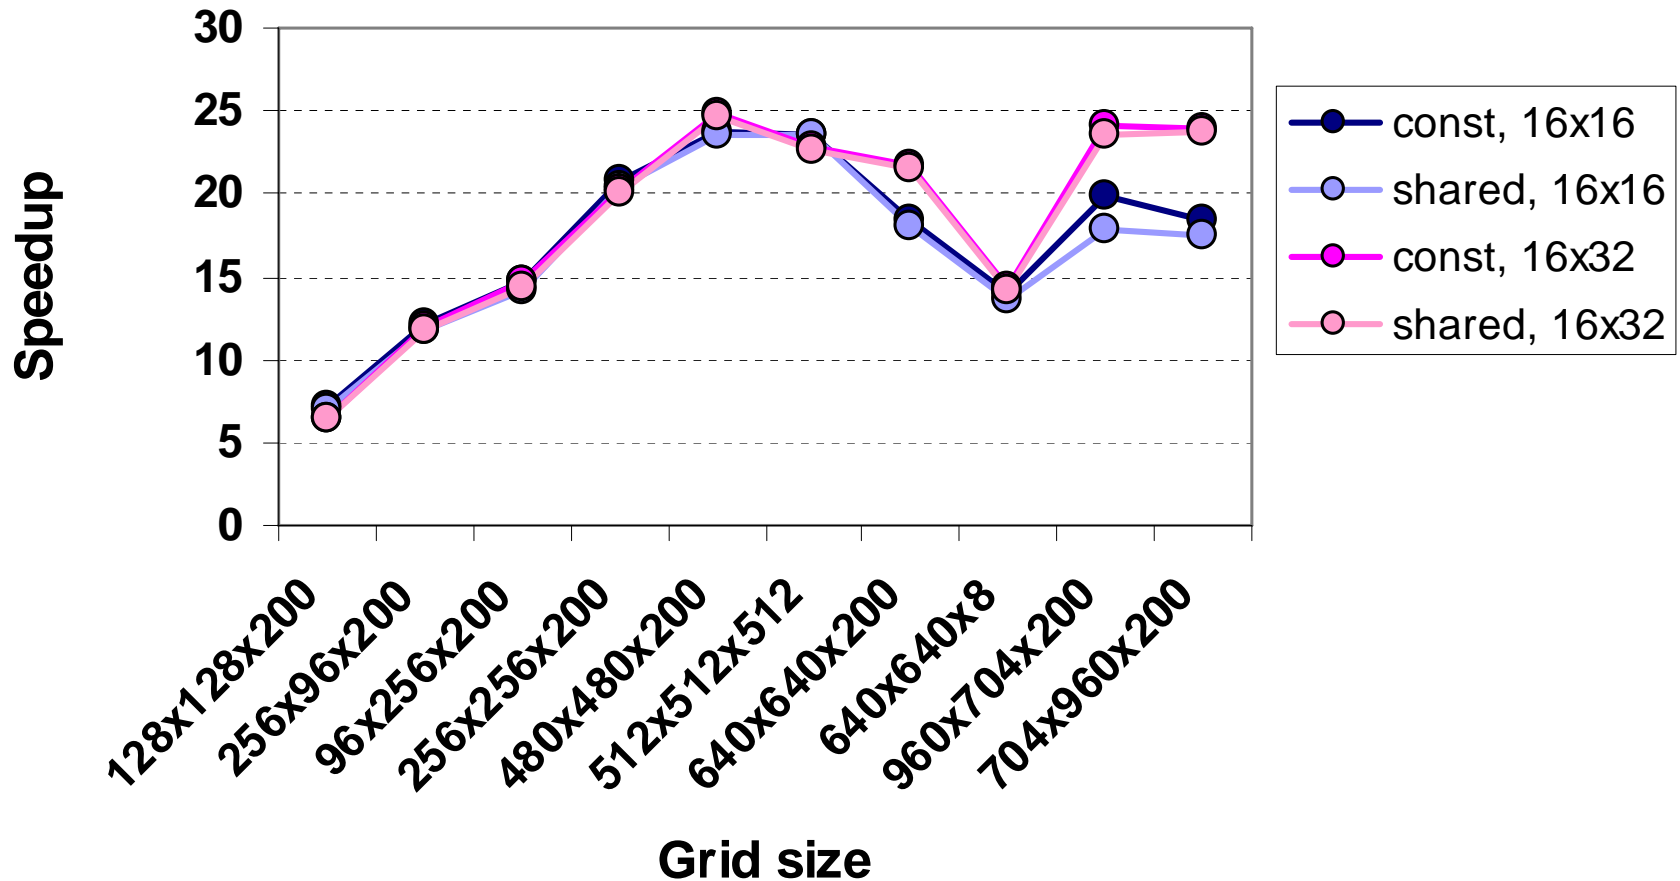


“Reverse-time” Imaging

Kernel performance



8th order, RTM



“Reverse-time” Imaging

Inter-GPU communication



- High frequency requires
 - Dense sampling
 - Large memory
 - Multiple GPUs
 - Halo exchange
 - Inter-GPU communication
- Device \leftrightarrow Host
 - Use pinned memory
 - PCIe bus predictably yields ~ 5 GB/s
 - ~ 10 % of kernel time
 - Easily hidden

“Reverse-time” Imaging

Inter-GPU communication

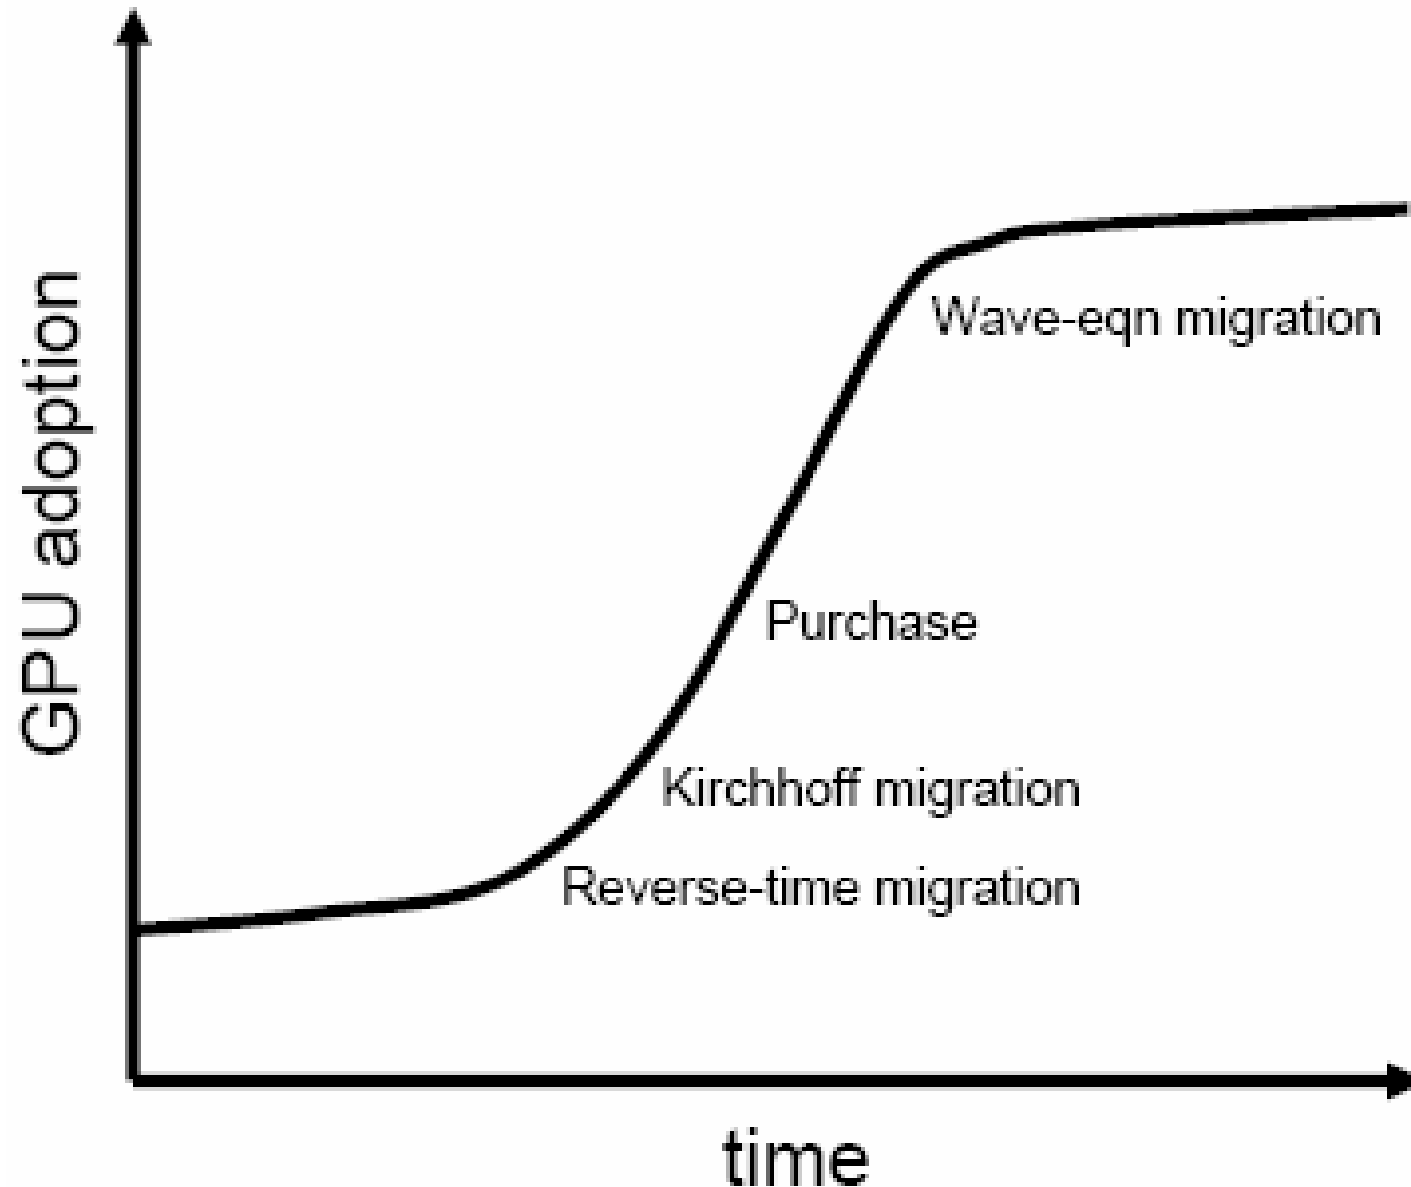


- CPU process \Leftrightarrow process
 - Currently using MPI
 - From legacy code
 - Performance variable
 - Comparable to kernel time
 - Solutions
 - OpenMP?
 - single controlling process?
- Node \Leftrightarrow node
 - Currently Gigabit Ethernet
 - Solution? Infiniband? 10-GigE?

- **Seismic imaging CUDA codes**
 - All 3 main codes written & verified
 - Two in production
 - One in production testing/optimization
 - All done with about two-man years of effort
 - Kernel speed-ups vary from 10 - 80 X on GT200
 - 456-GPU cluster out-performs 3000-CPU cluster
- **GPU cluster**
 - Jan 2008: bought 32-nodes (128 G80 GPUs)
 - Dec 2008: upgraded & expanded to 456 GT200s
 - Nov 2009: expanding to 1200 GT200s

Seismic Imaging

Summary





- **Code co-authors/collaborators**
 - Thomas Cullison (Hess & Colorado School of Mines)
 - Paulius Micikevicius (NVIDIA)
 - Igor Terentyev (Hess & Rice University)
- **Hess GPU systems**
 - Jeff Davis, Mac McCalla
- **NVIDIA support & management**
 - Ty Mckercher, Paul Holzhauer, Jeff Saunders, Philip Nenon
- **Hess management**
 - Jacques Leveille, Vic Forsyth, Jim Sherman